

Article

A Connected and Autonomous Vehicle Reference Architecture for Attack Surface Analysis

Carsten Maple * , Matthew Bradbury , Anh Tuan Le  and Kevin Ghirardello 

WMG, University of Warwick, Coventry, UK

* Corresponding Author: cm@warwick.ac.uk

Academic Editor: name

Version 14th November 2019 submitted to Appl. Sci.

1 **Featured Application:** The reference architecture presented is to be instantiated with different
2 components which is then used to analyse the attack surface of those components.

3 **Abstract:** Connected autonomous vehicles (CAVs) will be deployed over the next decade with
4 autonomous functionalities supported by new sensing and communication capabilities. Such
5 functionality exposes CAVs to new attacks that current vehicles will not face. To ensure the safety
6 and security of CAVs it is important to be able to identify the ways in which the system could be
7 attacked and to build defences against these attacks. One possible approach is to use reference
8 architectures to perform an attack surface analysis. Existing research has developed a variety
9 of reference architectures, but none for the specific purpose of attack surface analysis. Existing
10 approaches are either too simple for a sufficiently detailed modelling or require too many details to
11 be specified to easily analyse a CAV's attack surface. Therefore, we propose a reference architecture
12 using a hybrid Functional-Communication viewpoint for attack surface analysis of CAVs, including
13 the Devices, Edge, and Cloud systems CAVs interact with. Using two case studies, we demonstrate
14 how attack trees can be used to understand the attack surface of CAV systems.

15 **Keywords:** Connected Autonomous Vehicles; Reference Architecture; Attack Surface Analysis

16 1. Introduction

17 In recent years interest in deploying connected autonomous vehicles (CAVs) on real road networks
18 has been increasing [1]. In order to enable the applications that depend on connectivity [2] and
19 autonomy [3], vehicle computer systems are becoming more complex and there are becoming a greater
20 number of ways in which the vehicles can communicate with other devices, each other, nearby Edge
21 infrastructure, and the Cloud. Such changes in complexity [4], connectivity, and levels of autonomy
22 means that there are more ways in which a CAV can be attacked [5] and a successful breach carries
23 greater impact.

24 Due to the safety ramifications, it is important to protect the security of vehicles and the systems
25 they rely on. Security breaches could lead to vehicle theft, privacy leakage, or in the worst case lead
26 to injury or death of occupants. Analysing these security threats in isolation is insufficient since
27 vulnerabilities could be, and often are, exploited in combination to lead to escalated threats with the
28 potential for greater harm.

29 Reference architectures can be used to help understand and analyse complex systems, specifying
30 the entire system and any interactions. In addition to being a useful tool for analysis, a reference
31 architecture can be used to assist in performing attack surface analysis, for example, as part of the
32 system level analysis and design in SAE J3061 [6, Figure 7]. By using output from a threat modelling,
33 the identified goals, resources, capabilities, motivations and presence of an attacker can be used with

34 a reference architecture to help understand *how* an attack could be executed. However, a problem
35 with using existing reference architectures for attack surface identification and analysis is that they are
36 often either lack important details [7] in order to derive certain categories of attacks, or too complex [8]
37 for vehicle manufacturers and CAV system designers to feasibly use (which will be elaborated on
38 in [Subsection 2.1](#) and [Subsection 2.2](#)). This paper addresses these issues by proposing a hybrid
39 Functional-Communication viewpoint reference architecture for attack surface analysis. This reference
40 architecture aims to balance the complexity-completeness trade-off, such that the model is sufficient
41 complex to model a wide range of interactions but remains easy enough to practically use.

42 While many of the attacks against traditional vehicles could be modelled using this reference
43 architecture, we target L3–L5 autonomous vehicles (which are described in [Table 1](#)). These are the
44 new and emerging autonomous vehicles that are beginning to be deployed, and which will encounter
45 new threats compared to L0–L2 vehicles [9]. These new threats may try to manipulate input sensor
46 data [10] in order to affect how and where an autonomous vehicle drives, or may simply try to remotely
47 take control of the vehicle’s functions [11]. There is the potential for these attacks (and others [12])
48 to have a large impact due to the potential of leading to unsafe conditions for vehicle occupants and
49 pedestrians [13]. As the way in which vehicles are designed and operated is changing at a rapid
50 pace, this reference architecture aims to focus on the next 10 years [14] of autonomous vehicles and be
51 flexible to facilitate future changes.

52 To demonstrate the effectiveness of using this reference architecture to perform attack surface
53 analysis, we instantiate it with two different case studies. Using the interactions of components in
54 the reference architecture and goals identified from a threat modelling, attack surfaces are derived.
55 Performing the threat modelling to identify attacker goals, motivations, capabilities and resources is
56 out of the scope of this paper as the attack surface defines how these goals can be reached, but does not
57 aim to specify what these goals are. There exist many threat modelling approaches [6,15–18] that can
58 be used as input to the reference architecture. In the first example of valet parking, the attacks against a
59 vehicle parking itself in an autonomous car park are investigated. In the second example, a real world
60 attack against Tesla vehicles is used to highlight the need to consider the Edge infrastructure in the
61 security of CAVs.

62 We make the following contributions in this paper:

- 63 1. A reference architecture made up of 4 sub-architectures: CAVs, Devices & Peripherals, the Edge,
64 and the Cloud formed of a hybrid Functional and Communication viewpoint.
- 65 2. A methodology to use the reference architecture to synthesis the attack surface in the form of
66 attack trees.
- 67 3. Two case studies to demonstrate the applications of attack surface and attack tree analysis in
68 deepening the security knowledge of the system.

69 The remainder of this paper is organised as follows. In [Section 2](#) we present relevant related
70 work, including examining existing vehicular reference architectures. [Section 3](#) describes our proposed
71 reference architecture, its components and relevant attack surfaces, and in [Section 4](#) describes the
72 methodology for using the attack surface; including using attack trees as a method to perform the
73 analysis. In [Section 5](#) two case studies of example applications are presented as instantiations of our
74 reference architecture. The implications of the reference architecture is discussed in [Section 6](#); and
75 future work is presented in [Section 7](#), before the paper concludes with [Section 8](#).

76 2. Related Work

77 There has been much work conducted on analysing the threats that an autonomous vehicle
78 will face [7,11,21–23]. The issue with existing work on threat analysis is that they did not consider a
79 comprehensive ranges of components (i.e. CAV, devices, Edge, Cloud) that form the potential CAV
80 operational contexts. This means that threats which use a combination of attacks against different
81 components in specific orders can be missed. Reference architectures have been developed to aid in

Table 1. Levels of Vehicular Autonomy [19] from no autonomy (where the driver is in full control of the vehicle) at level 0 to level 5 where the vehicle is in full control.

Level	Name	Description	Example
0	None	The human driver is in full control.	Anti-lock Braking System
1	Driver Assistance	The human driver is assisted by a driver assistance system of steering or acceleration/deceleration using information about the driving environment. The human performs all other tasks.	Cruise Control
2	Partial	The human driver is assisted by one or more driver assistance systems of both steering and acceleration/deceleration using information about the driving environment. The human performs all other tasks.	Lane Centring
3	Conditional	The autonomous vehicle controls all aspects of driving, with the expectation that the human driver will respond appropriately to a request to intervene.	Traffic Jam Chauffeur [20]
4	High	The autonomous vehicle controls all aspects of driving, even if a human driver does not respond appropriately to a request to intervene.	Driverless Valet Parking [20]
5	Full	The autonomous vehicle is in full control and no human input related to driving is expected.	Driverless Cars

82 the design of products and services for autonomous vehicles, but have seldom been used to provide
83 a wider view of composite threats. Those reference architectures that do exist, can suffer from being
84 too broad, or are insufficiently detailed, for attack surface analysis. When too broad, they require
85 specifying less pertinent details as part of the model, which detracts from performing the attack surface
86 analysis. When insufficiently detailed, there are threats that cannot be analysed using the reference
87 architecture. The remainder of this section will present related work on reference architectures used to
88 model autonomous vehicle systems.

89 2.1. Reference Architectures

90 In order to better analyse how a system is structured *reference architectures* are used as an
91 abstract way of specifying a system. A reference architecture is an approach to model a system
92 and provide a consistent and standardised way to describe that system. This common model should
93 be created such that it is able to describe a broad range of scenarios that the system can be used in.
94 Reference architectures allow modularisation of a system into components and interfaces between
95 these components to be defined. These features can be used to assist with system development in a
96 scalable way (e.g., by multiple organisations [24]) and also facilitate testing of the system.

97 This paper will develop a reference architecture specifically for assisting in the attack surface
98 identification and analysis in CAVs. We will present our reference architecture in the next section,
99 but provide here an overview of existing reference architectures. This related work guides our own
100 development, and assists us in identifying the shortcomings of existing schemes that are discussed in
101 the next section.

102 A common feature of reference architectures is to decompose the system they are modelling into
103 multiple viewpoints and then specify those viewpoints in detail. There are several different viewpoints
104 that reference architectures can present, including:

- 105 • Functional: how the components work and what their tasks are
- 106 • Communication: how the components interact
- 107 • Implementation: how the components are implemented
- 108 • Enterprise: the relation between organisations and users
- 109 • Usage: concerns of expected usage of the system

- 110 • Information: the types of information handled by the system [25]
- 111 • Physical: the physical objects in the system and their connections

112 Of these viewpoints, presenting the Functional, Communication, and Implementation tend to be the
113 most common as they cover what the system does and how the system interacts with itself and other
114 systems. When developing a reference architecture, it is important to develop only the viewpoints
115 necessary to describe the system to prevent a user of the reference architecture needing to provide
116 additional unnecessary information.

117 2.1.1. Non-CAV Reference Architectures

118 Before exploring the existing CAV reference architecture it is useful to examine reference
119 architectures for different fields. In doing so they raise interesting ideas for ways in which CAV
120 reference architectures can be improved.

121 A common architectural framework for the development of interoperable industrial internet
122 systems was presented in [26]. The Industrial Internet Reference Architecture (IIRA) is divided
123 into four viewpoints, namely Business, Usage, Functional and Implementation. While the last two
124 viewpoints are of utmost importance in the identification of a system's threats and vulnerabilities, as
125 they are concerned with a system's functional requirements, interdependencies, and technological
126 implementations. The IIRA also describes a system's business objectives and expected usage, both of
127 which go beyond the scope of a reference architecture for attack surface analysis.

128 A Smart Grid Reference Architecture was developed in [27] which uses Business, Functional,
129 Information, and Communication viewpoints. Explicit considerations of *information security*, are
130 included (i.e., confidentiality, privilege escalation), however, the methodology of how to perform
131 a security analysis of the system is not described. The systems described are complex and include
132 many implementation details, including the scenario a component is operating in and what actions
133 the component is involved with. From a security analysis perspective the reference architecture
134 could be simplified (e.g., by removing business cases) to reduce the scope for which cyber security
135 needs to be considered. This means that while the reference architecture states that it is useful for a
136 cyber security analysis, due to it describing aspects of a Smart Grid which do not have cyber security
137 considerations, performing a cyber security analysis is difficult. The conclusions from this are that
138 reference architectures for cyber security analysis, should focus on the aspects of the system for which
139 cyber security is relevant.

140 2.1.2. CAV and ITS Reference Architectures

141 A functional reference architecture for autonomous driving was introduced in [28], which
142 provided a foundation for considering the functionality of an autonomous vehicle irrespective of
143 its implementation. There are close relations between functional safety and security analysis in
144 the automotive domain. The functional safety analysis relies on information taken from hazard
145 identification, which can be influenced by security aspects such as the communication between the
146 components or access to assets. On the other hand, the implemented countermeasures to address
147 functional safety can determine the security level of the system. As a result, there are certain attempts
148 to integrate security into (functional) safety analysis in CAV, such as SAE J3061 [6]. However, there is
149 insufficient focus on CAV interactions to support using this model for an attack surface analysis. This
150 is because the approach focuses on the vehicle only, and does not consider interactions with RSUs,
151 other vehicles, the Internet, and other devices.

152 In [7] a security-focused risk assessment was performed for autonomous driving (AD). To achieve
153 this the authors defined a reference architecture by synthesising from multiple academic and industrial
154 AD sources to model select AD applications. The model was instantiated for different selected
155 applications of interest and a risk assessment of the identified threats was performed. The authors
156 note that their work does not attempt to perform an exhaustive specification of threats, but to provide
157 ways to specify the system to aid in deriving the threats. The reference architecture and analysis of it

performed in our work is similar to this paper, however, we argue that certain details are lacking from this model which prevents a sufficiently in-depth analysis of the attack surface.

A reference architecture for ITS infrastructure that focuses on business and organisational aspects of the system was presented in [29]. While the paper does not discuss technical considerations of an ITS system, the organisational aspects highlight certain areas of interaction which are of interest from a security perspective. One issue that was highlighted was that heterogeneous systems had trouble interacting due to different implementations by different suppliers. An adaptor was required to allow these systems to interact, which would be a component of the attack surface. The reference architecture raises the importance of service collaboration, for example, parking and guidance services will need to collaborate to ensure a car is not directed to a full car park. The interactions between these services will also form part of the attack surface.

A detailed and comprehensive reference architecture for cooperative and intelligent transport was developed in [8]. There are three components to this architecture, (i) Architecture Reference for Cooperative and Intelligent Transportation (ARC-IT), (ii) Regional Architecture Development for Intelligent Transportation (RAD-IT), and (iii) Systems Engineering Tool for Intelligent Transportation (SET-IT). RAD-IT focuses on tools for regional ITS architectures and SET-IT focuses on assisting with developing “architectures for pilots, test beds and early deployments”. The key component is ARC-IT when it is used to specify a Functional viewpoint¹ and Communication viewpoint². The architecture is designed to be comprehensive, which is a benefit as the architecture can be used to specify interactions in detail. However, the additional detail adds additional complexity that makes the tool harder to use. There is need for a simpler model that can be easily analysed.

The CARMA project [30], which aims to investigate the distribution of the autonomous control functions throughout an ITS defines a three tiered architecture in terms of the CARMA CORE, CARMA EDGE and VEHICLE. The CARMA CORE layer acts as in a supervisory role of the distributed vehicle control functions (such as mission planning of the end-to-end vehicle trip). The majority of mid-level controls, such as improving the calculation of reference signals for vehicle control, are implemented in the CARMA EDGE. However, some of these mid-level controls are implemented in the VEHICLE layer. The CARMA system presents a model of a complex autonomous system that introduce a number of security concerns and challenges [31]. A reference architecture could be used to achieve an understanding of the attack surface thereby allowing a more holistic threat assessment.

ITS reference architectures have also been developed for other regions, such as Holland [32], the USA [33], and Europe [34]. However, these architectures suffer from the same problem that ARC-IT does, that they are intended to be very general and cover a wide range of considerations of intelligent transport systems. This lack of focus reduces their usability to undertake an attack surface analysis.

2.2. Requirements for Attack Surface Analysis

The extant reference architectures for CAVs variously consider analysis (of attacks and of risk), viewpoints, and features (autonomous vehicles, devices, edge and cloud). The reason that these architectures have different characteristics is that they serve different purposes. When analysing an attack surface, not all of the characteristics are required, indeed some are undesirable as they may be too detailed and complex, and as such not be effective for easy identification of the surface and associated threats. To be most effective, a reference architecture needs to have the essential characteristics and no more. For example, [8] considers the widest range of viewpoints, but this can hamper the security analysis. One example of this is that the information flow of the system is described in the Physical Viewpoint using entities from the Enterprise View. These information flows are also described in the Communications viewpoint. This repetition is helpful for system design within a single viewpoint, but

¹ <https://local.iteris.com/arc-it/html/viewpoints/functional.html>

² <https://local.iteris.com/arc-it/html/viewpoints/communications.html>

Table 2. Summary of CAV Reference Architectures where the purpose, viewpoints used and the components are identified with a ✓ if included or an ✗ if not included.

Reference Architecture	Analysis		Viewpoints						Considers				
	Attack	Risk	Functional	Communication	Implementation	Enterprise	Usage	Information	Physical	CAV	Devices	Edge	Cloud
Behere and Törngren [28]	✗	✗	✓	✗	✗	✓	✗	✗	✗	✓	✗	✗	✗
Osório <i>et al.</i> [29]	✗	✗	✗	✗	✗	✓	✗	✗	✗	✗	✗	✗	✗
Dominic <i>et al.</i> [7]	✓	✓	✓	✓	✗	✗	✗	✗	✗	✓	✓	✓	✗
The Architecture Team [8]	✓	✗	✓	✓	✗	✓	✓	✗	✗	✓	✓	✓	✓
Passchier and van Sambeek [32]	✗	✗	✓	✓	✗	✗	✗	✗	✓	✓	✗	✓	✓
Heise [33]	✗	✗	✓	✗	✗	✓	✗	✓	✓	✗	✗	✓	✗
Begoña <i>et al.</i> [34]	✗	✗	✓	✓	✗	✗	✗	✓	✓	✗	✗	✓	✓
This Work	✓	✗	✓	✓	~	✗	✗	✗	✗	✓	✓	✓	✓

not security analysis across multiple viewpoints; a more focused reference architecture can simplify the process of performing a cyber security attack surface analysis.

The minimal viewpoints required for a cyber security attack surface analysis are Functional and Communications, as is it necessary to know what the system does and how it interacts. This allows what actions an adversary can perform and how an adversary's interactions with the system can produce the attack. Other components are necessary for other systems, for example, the Physical viewpoint is required to investigate cyber-physical attacks. Other viewpoints, such as the Implementation viewpoint are important to analyse attacks against specific systems. But to perform a more general attack surface analysis, the Functional viewpoint is sufficient. Other viewpoints (e.g., Enterprise and Usage) are useful in considering different types of security such as security management. Therefore, the Functional and Communications viewpoint can be focused on when performing a cyber security attack surface analysis.

A comparison of the existing and proposed reference architectures is provided in Table 2. Features that the reference architecture includes is indicated with a ✓, and features that are not included are indicated with a ✗, the following features are shown: (i) purpose of the reference architecture (Analysis), (ii) the viewpoints used (Viewpoints), and (iii) the areas the reference architectures consider (Considers). Our work partially considers the Implementation viewpoint as it can be implemented as a *virtual* component and is thus marked with a ~. Some of the existing reference architectures fail to focus on the wide range of interactions that a CAV could be involved with. Most reference architectures include Edge devices such as RSUs, but do not considering the wider range of interactions between CAVs, Devices & Peripherals, the Edge and the Cloud. Without considering all of these interactions, it will be impossible to analyse many current and emerging attacks, so a new reference architecture needs to specify these interactions.

2.3. Summary

There are many threats that have been identified for CAVs and there have been several reference architectures developed to analyse the attack surface of CAVs. However, the reference architectures tend to either be too broad and consider aspects of an ITS that do not need to be specified when considering the attack surface of CAVs, or lack sufficient detail to analyse certain types of threats. In the next section we will present a reference architecture formed of a hybrid functional-communication viewpoint to address the lack of reference architectures that balance ease of use with being sufficiently detailed.

234 3. A CAV Reference Architecture: Components and Related Attack Surfaces

235 The reference architecture presented in this work uses the Functional and Communication
 236 viewpoints combined into a single hybrid viewpoint. These are the minimal two viewpoints needed,
 237 as a threat agent would need to know what the CAV does and how the CAV can be interacted
 238 with to attack it. However, the Implementation is also an important viewpoint (as will be shown in
 239 [Section 5](#)), because a threat actor can take advantage of vulnerabilities in the implementation of a
 240 component. To resolve this in our reference architecture, the implementation can be considered as
 241 part of a functional component, or as a *virtual* functional component that exists and interacts with all
 242 components. Important virtual components that might exist include the Operating System and the
 243 hardware that the software is executing on (e.g., Electronic Control Units (ECUs)). The users of the
 244 system are considered when identifying the scenarios of interest in which the reference architecture
 245 will be instantiated with concrete components. Finally, how users and organisations interact may lead
 246 to security issues (e.g., resetting a password), but as these threats do not specifically relate to CAVs
 247 they are out of the scope of this paper.

248 The four sub-architectures that are presented were designed by identifying key components
 249 within CAVs and the ways in which they will interact. The sub-architectures for CAVs and Devices &
 250 Peripherals are presented in [Figure 1](#). The two sub-architectures for the Edge and the Cloud are shown
 251 in [Figures 2](#) and [3](#) respectively. These architectures are composed of various abstract components which
 252 need to be instantiated with concrete implementations to undertake an analysis of the architecture. For
 253 example, the *Sensors* component could be instantiated with GPS, LIDAR, tire pressure, and temperature
 254 sensors. These components should be instantiated with the desired concrete implementations that are
 255 required for a specific application. When analysing different applications, the reference architecture
 256 will be instantiated with a different set of components.

257 3.1. CAV Reference Architecture

258 The first of three reference sub-architectures is shown in [Figure 1](#), and it specifies the abstract
 259 components for CAVs and the devices & peripherals that interact with the CAV. Certain components
 260 are not included in the diagram as they are implementation details. For example, how the components
 261 interact (internal communications, usually via the Controller Area Network (CAN)), how the
 262 components are implemented (usually as an ECU), or what operating system is used. These
 263 components are important to consider when analysing attacks, but they do not form the high level
 264 functionality of the system. For example, the telematics control unit subject to research in [\[22\]](#) contains
 265 multiple functional components in a single physical component. The remainder of this section will
 266 describe the components present in the architecture.

267 3.1.1. Wireless Communications

268 Vehicles are currently or expected to be equipped with
 269 multiple antennas in order to communicate over different
 270 wireless protocols. This includes antennas for (i) receiving audio
 271 over AM, FM or DAB radio, (ii) receiving and transmitting
 272 IEEE 802.11 WiFi, (iii) bidirectional V2X communications over
 273 IEEE 802.11p, and (iv) bidirectional cellular antennas (such
 274 as 4G). It may also be the case that Internet of Things (IoT)
 275 technologies such as IEEE 802.15.4 or ZigBee are included
 276 to facilitate interoperability with IoT networks. Many of the
 277 systems in the CAV will interact with the communications due
 278 to the need to coordinate with nearby vehicles, or to provide
 279 services to the vehicle's users. As communications are the
 280 primary way in which vehicles will exchange information, they

Example Attacks
<ul style="list-style-type: none"> • DoS V2X communications [35] • Eavesdrop • Replay • MiTM Intercept • Incorrect handling of malicious packets (e.g., DAB [36]) leading to RCE • Context information leakage (e.g. location, identity [37]) • Sybil Attacks [38] • Colluding to defeat agreement protocols [39] • Wormhole (Relay) Attack [40]




Figure 1. CAV and Devices & Peripherals Reference Architectures (Hybrid Functional-Interaction viewpoint)

281 will be the avenue through which most attacks are launched. These attacks may try to compromise or
 282 interfere the way in which packets are communicated, or compromise the components to which the
 283 packets are forwarded.

284 3.1.2. Physical Inputs and Output

285 Physical inputs and outputs that are contained within a
 286 vehicle include ports such as USB, OBD-II, audio connections,
 287 and others. Exploiting these ports is typically harder for an
 288 attacker as they would usually require physical access to the
 289 vehicle, however, due to the presence of additional devices that
 290 connect to these ports there are ways in which attacks can be
 291 performed via a remote connection.

292 With the presence of a USB port (depending on the protocol with which the hardware interprets
 293 the data), there is the possibility for an adversary to gain access to the vehicle's internal network [42].
 294 Malicious USB sticks could be given out to people loaded with music or videos for free, with the
 295 intent of being plugged into the vehicle. When plugged in, malware could attempt to access the
 296 internal vehicle's CAN bus. Another approach is to fool users into connecting a device that resembles
 297 a USB stick but can repeatedly deliver a high voltage discharge that would destroy a vehicle's internal
 298 electronics [41].

Example Attacks

- Cause electrical damage [41]
- Install malicious software (e.g., by firmware updates on CDs or USB sticks) [11]

3.1.3. Internal Communications (Virtual)

As well as a communication system that allows a CAV to communicate with external devices (such as RSUs or 3rd party vehicles), they also have an internal communication system such as a CAN bus. This is used to connect the multiple components that form an implementation of the functions specified in the architecture. This communication network is not explicitly specified as a component, as it is implicit due to the components interacting. Vehicles may also use a different internal interaction (such as Ethernet) in the future and by under-specifying this implementation detail, the reference architecture is more generic.

Attacking this internal communication network can be performed by a direct connection to it, for example, via an OBD-II port. Alternatively, an attacker can gain access to this internal network via vulnerabilities in the components that connect to it. Once these components are compromised an attacker will have access to eavesdrop on messages sent [43], or the ability to inject malicious messages [42–44]. With access to the internal network of a vehicle many functional aspects of the vehicle can be controlled, including: the radio, instrument panel, the vehicle’s body, engine, brakes, HVAC, and others [21]. A solution to these issues is to use encryption and authentication of messages [45], however, vehicles currently on the road act as if the CAN bus is a walled garden and do not attempt to encrypt or authenticate messages sent on the bus.

Example Attacks
<ul style="list-style-type: none"> • Send crafted packets [11,42–44] • DoS [43] • Eavesdrop [43]

3.1.4. Sensors (Internal and External)

Sensors are a key component of CAV systems. The vehicles will rely on their input to build a model of the world. Example sensors include: (i) Global Navigation Satellite System (GNSS) to be aware of a vehicle’s position, (ii) wheel rotation sensors to be aware of velocity, (iii) LIDAR to be aware of the relative position of other vehicles, (iv) parking cameras to assist drivers, and a variety of other sensors such as temperature, humidity and light. Sensors may also observe information *passively* about the occupants in the vehicle. As sensors are a way for the vehicle to obtain the state of the environment around it, if that data can be maliciously manipulated, then the vehicle may make incorrect decisions based on the manipulated data. Alternately, an adversary may attempt to eliminate the vehicle being able to use certain sensors, such as by jamming GNSS signals or producing too much LIDAR interference for the data to be useful [10]. Another approach may be for an adversary to place additional sensors on the vehicle exterior, or to subject the sensors to physical manipulation.

In certain systems, the vehicle’s sensors may wirelessly communicate their data to the car (such as when monitoring tire pressure [46]). Most sensors are expected to be hardwired to the system due to high reliability requirements. Wireless sensors pose a greater security threat as there is a larger attack surface for an adversary to take advantage of. For example, the Tire Pressure Monitoring System (TPMS) leaks identity information about the vehicle by including unencrypted identifiers in the packets it sends. Due to the lack of authentication and validation, the system also is vulnerable to spoofing and replay attacks, where the vehicle could easily be fooled into believing the tire is flat even if it was not.

Example Attacks
<ul style="list-style-type: none"> • Induce misleading readings (Spoof, Replay, Delay) [10] • Blind, Jam [10] • Tamper (Disable, Replace)

3.1.5. Data Storage

Vehicles will need to store data, including (i) the firmware and software used to run the car, (ii) maps and navigation information, (iii) music and videos for the entertainment system, and other information necessary for different use cases. This data will not be stored in a central location on the vehicle and will be stored in multiple locations. Data storage should also

Example Attacks
<ul style="list-style-type: none"> • Violate Integrity (manipulate data) • Violate Confidentiality (extract data) • Violate Availability (delete data) • Violate Non-repudiation (delete logs) • Remote firmware update [22]

346 be segregated based on the purpose for the data. For example,
 347 music and video should not be stored in the same location as the
 348 vehicle's software, but implementation details may mean that
 349 this is not the case. Not all data will be stored locally, some will be present in the Cloud and only
 350 requested when required. Other data may be stored in the Edge, or even in other vehicles on the road.

351 3.1.6. Data Analysis

352 To make sense of the data obtained from external sources
 353 (such as the sensors) and the data stored locally in the vehicle,
 354 some sort of analysis will need to be performed on it. This
 355 analysis may use simple conditions to trigger actuators (e.g.,
 356 if temperate rises above a threshold, then turn on the air
 357 conditioning), but more complicated techniques such as machine
 358 learning models will also be used. These machine learning
 359 models will be prevalent in CAVs due to the need for autonomy.

Example Attacks
<ul style="list-style-type: none"> • Induce bad analysis (e.g., adversarial ML [47]) • Obtain analysis • Malicious input to put analysis into infinite loop (DoS)

360 Localisation

361 One of the key pieces of knowledge for an autonomous vehicle is its location. Information such as
 362 from GNSS can be used to provide a fairly accurate location [48] as long as the vehicle is in an open
 363 area with few buildings blocking satellite signals. Other approaches such as dead reckoning are used
 364 to calculate the vehicle's current position based on a previously known position, the vehicle's speed,
 365 heading, and the travel time.

366 Object Identification

367 As part of autonomous driving it will be necessary for the CAV to be able to identify objects.
 368 These objects will include people, obstacles, road signs, and many other objects. Machine learning
 369 based algorithms will be used to perform visual identification. However, using machine learning can
 370 open the vehicle up to being attacked in new ways. One example is adversarial machine learning,
 371 where input manipulation can lead to unexpected results. For example, in [49] 3D printed objects
 372 were crafted to be misclassified by an object detection model. In one case a turtle was detected as
 373 a gun, such a detection could lead to unexpected behaviour in the vehicle. Alternate issues might
 374 include the vehicle failing to recognise another vehicle, such as when a Tesla was involved with a fatal
 375 accident when it attempted to drive under a truck [50]. An adversary manipulating the data provided
 376 to sensors, may affect the actions vehicles take.

377 Sensor Fusion

378 To improve accuracy of sensor input the data provided from sensors is usually fused, such as via
 379 a Kalman Filter [51]. By doing so the quality of the fused data should be higher than the individually
 380 sensed data. However, if manipulated sensor data is used then the fusion approach could produce
 381 less accurate or even inaccurate results [52]. In [53] spoofing sensor data was used to control a UAV,
 382 with the technique possibly extendable to other autonomous vehicles. Therefore, the sensor fusion
 383 technique needs to be aware of how to handle data provided by an attacker, such that it does not lead
 384 to incorrect actuations.

385 Action Engine

386 Once an autonomous vehicle has both determined its location and the road objects surrounding
 387 it, it may call on the Action Engine sub-module to decide what it must do next. Possible actions to be
 388 taken include interactions with other connected vehicles on the road and both short and long term
 389 driving decisions. RSUs or the Cloud, on the other hand, make use of the Action Engine to ensure

390 that the vehicle control and planning systems are correct and safe, and to ensure that multiple vehicles
 391 on the road at the same time coordinate and are managed to move people and packages to their
 392 destinations in the most effective way.

393 3.1.7. Energy System

394 The energy system both supplies energy (in the form of
 395 electricity) to the systems within a CAV and is also capable of
 396 being supplied with energy. Energy can be supplied back to the
 397 batteries through the use of regenerative breaking, solar panels,
 398 recharging cables, and other sources. The energy system is also
 399 tasked with maintaining the vehicle's batteries to ensure power is safely drawn from them. If the
 400 energy system is compromised then unsafe usages of electricity may follow which could lead to
 401 damage to the vehicle.

Example Attacks

- Overcharge battery to damage it
 - Drain power
-

402 3.1.8. Actuation

403 This module contains any components that can perform an
 404 action with an impact on the physical world. This may include,
 405 applying the brakes, changing wheel speed, changing the angle
 406 the wheel is pointed in, operating the air conditioning, lowering
 407 or raising windows, locking and unlocking car doors, and others. If an adversary is not attempting
 408 to gain information about the vehicle or passengers, then actuating components are likely to be the
 409 key target. For example, an attacker may attempt to compromise a large number of vehicles in order
 410 to provide Theft as a Service (TaaS) [54]. Rather than steal cars, the thief will install malware on as
 411 many vehicles as possible. Then, when there is demand for a particular car the malware can give the
 412 thief access to the vehicle. The adversary who installed the malware may not even need to active the
 413 malware themselves, as they could provide a crafted key to the intended buyer.

Example Attacks

- Disable
-

414 3.1.9. Monitoring and Logging

415 Monitoring and logging are important aspects for CAVs
 416 in a number of scenarios, including: verifying that vehicles are
 417 functioning correctly, analysing past decisions made, and will be
 418 used to manage maintenance schedules. For example, if a CAV is
 419 in an collision the vehicle will need the ability to explain why it made the decisions before the collision.
 420 If an adversary is capable of accessing the diagnostics unit then it may rewrite decision making history,
 421 preventing reliable auditing.

Example Attacks

- No longer forensically valid
 - Extract data
-

422 3.1.10. Infotainment

423 The infotainment system is used to manage the
 424 entertainment system within a vehicle (such as audio/video
 425 systems) and information systems (such as maps and navigation,
 426 phone, and car status). Infotainment systems are also likely to
 427 contain a web browser to facilitate access to the internet for both
 428 entertainment and information. An issue with infotainment
 429 systems is that they may process data from untrusted sources. If the data is maliciously crafted to take
 430 advantage of vulnerabilities in the system, then an attacker may be able to remotely execute arbitrary
 431 code.

Example Attacks

- Arbitrary code execution (via browser) [43]
 - Arbitrary code execution (via crafted audio/video files)
-

432 3.1.11. Human-Machine Interface

433 A Human-Machine Interface (HMI) is any device or
 434 software which allows a person to *actively* interact with a
 435 machine. A passive observation of the occupants would be
 436 performed by the Sensors component. In vehicles HMI includes
 437 critical systems such as the steering wheel, accelerator pedal, break pedal, and gear controls. Less
 438 critical system include the controls on the dashboard and feedback mechanisms. An attacker may
 439 attempt to intercept the signals from the HMI to prevent the vehicle doing something other than
 440 requested. Alternatively, the attacker might use the HMIs to report statuses that are incorrect to
 441 attempt to get the driver or passengers to perform certain actions. For example, the adversary may
 442 turn on engine safety warnings (when there is no problem) to cause the driver to stop the car. The
 443 attacker could then use this opportunity to steal the vehicle, or perform other attacks, such as attaching
 444 a tracking sensor.

445 Note that HMI does not communicate directly with the actuators. There will need to be some data
 446 analysis performed that potentially adjusts the action performed. For example, an anti-lock breaking
 447 system would not always actuate the brakes in the way the driver requests.

448 3.2. Devices and Peripherals Reference Architecture

449 Vehicles may have a number of peripherals that interact with each other. Some examples of the
 450 kinds of devices and peripherals that may be present and in use are: (i) Car Keys, (ii) Smart phones, (iii)
 451 MP3 players, (iv) Bluetooth devices, (v) 3rd Party Navigation Systems, (vi) Dashcams, (vii) Portable
 452 games consoles, and others. These devices could either interact with the vehicle or simply be present
 453 within the vehicle. Some of these interactions may be relatively simple, such as accessing the vehicle's
 454 WiFi in order to connect to the internet via a cellular connection. Others may involve accessing the
 455 vehicle's storage, actuating the infotainment system, or controlling other aspects of the vehicle. These
 456 peripherals are additional vectors that attackers can take advantage of to attack the system. This may
 457 be by loading the device with malware to gain control [55], or interacting with the context of the
 458 inter-device communication [56].

459 It is also the case that some of these interactions may be unintended. For example, a passenger
 460 leaving their phone in an automated taxi may leak the journey history of the taxi if it is running a
 461 phone tracking service. This sort of leak could also be caused by an attacker intentionally attaching
 462 such a device to the vehicle.

463 3.2.1. Applications

464 One of the key features of certain devices (such as
 465 smartphone) are the ability to run applications on it. Some
 466 vehicle manufacturers (such as Volkswagen [58]) are creating
 467 mobile apps that obtain information from the car or allow the
 468 app to control certain features (such as the infotainment system).
 469 If the phone is compromised then that malware may be able to
 470 affect the vehicle's systems via the app. The attacker may be able to leak data about the car, gain an
 471 internal vector to the vehicle's systems, or use the phone's connection to the cloud to attack the vehicle.

472 3.2.2. Sensors

473 The devices within a vehicle may have their own sensors
 474 that reveal information about the state of the environment inside
 475 the vehicle, or about the vehicle itself. An adversary may wish
 476 to take advantage of these sensors to gain knowledge about
 477 the vehicle, which could be potentially useful in escalating the
 478 severity of other attacks.

Example Attacks

- Spoofing vehicle status
 - Intercept commands
-

Example Attacks

- Location tracking via sensor data (e.g., magnetometer [57])
 - Data harvesting
 - Become internal attack vector for remote adversary
 - Malicious smartphone app interfering with CAN bus [55]
-

Example Attacks

- Blind, Jam
 - Induce misleading readings (Spoof, Replay, Delay)
-




Figure 2. Edge Reference Architecture (Hybrid Functional-Interaction viewpoint)

479 3.2.3. Wireless Communications

480 The devices present in a vehicle are expected to
 481 communicate wirelessly. This may be with the cellular network,
 482 directly with the vehicle, or possibly with other devices in the
 483 vehicle. One example, is that vehicular privacy may be leaked
 484 due to the presence of devices in the vehicle. For example, WiFi
 485 devices will broadcast their MAC addresses periodically when
 486 looking for a device to connect to [59]. Bluetooth devices will also
 487 beacon their MAC address in order to find devices to connect
 488 to [56]. Both reveal identity information that could be used to
 489 track people in vehicles.

Example Attacks
<ul style="list-style-type: none"> • Relay Attack (Car Key Signal [40]) • Replay attack (e.g., unlock car using recorded signal) • Wireless protocols leak identity information about owner [59] • Facilitates tracking of person and vehicle [56]

490 3.3. Edge Reference Architecture

491 The Edge reference architecture specifies the interactions of
 492 components that occur between operations of the vehicle and
 493 the operations of the Cloud. This may include devices used to
 494 access a WAN (such as cellular base stations or WiFi hotspots).

Example Attacks
<ul style="list-style-type: none"> • Modify hardware (Tamper) • Disable hardware

495 Edge devices must include some functionality that does not occur remotely but occurs close to where
 496 the vehicle is operating or at the boundary between the vehicle and the cloud. There is a wide
 497 range of scenarios that could be considered in the Edge reference architecture. The main example are
 498 Road-Side Units (RSUs) which are computing devices placed along road networks to support CAVs
 499 travelling along the roads. These devices will communicate with autonomous vehicles to assist their
 500 autonomous activities. Alternate pieces of infrastructure could also be considered as part of the Edge.
 501 For example, internet connected traffic lights, smart parking garages, and others, may need to interact
 502 with autonomous vehicles and actuate components to facilitate autonomous driving.

503 Certain components have been previously described (e.g., Sensors, Data Analysis) and will not be
 504 repeated as part of the Edge sub-architecture. Some components previously described will be repeated
 505 due to differences with the previous sub-architectures.

506 3.3.1. Communication

507 Communication on the Edge has additional capabilities
 508 compared with CAVs and the Devices & Peripherals within
 509 them, as the Edge can be physically connected to a wide area
 510 network (WAN) rather than just wirelessly connected. Such
 511 physical connections might be provided by high bandwidth fibre, Ethernet and other communication
 512 approaches that require a physical medium. However, Edge nodes will still need to have wireless
 513 communication in order to facilitate V2I communication. This communication will include the
 514 technologies specified in vehicles to facilitate Dedicated Short-Range Communications (DSRC) (e.g.,
 515 IEEE 802.11p and/or C-V2X). Other technologies might include non-vehicular specific cellular
 516 communications, WiFi and protocols to interact with IoT systems (e.g., IEEE 802.15.4).

Example Attacks
<ul style="list-style-type: none"> • Edge Emulation [60] • DoS

517 3.3.2. Data Storage

518 Data storage at the Edge will typically be centralised in
 519 each device as a single piece of hardware. As the Edge is
 520 susceptible to tampering it is important to ensure precautions
 521 such as encrypting the entire disk is used to prevent a threat
 522 actor from removing, reading from, and then replacing the disk.

Example Attacks
<ul style="list-style-type: none"> • Violate Integrity (manipulate data) • Violate Confidentiality (extract data) • Violate Availability (delete data)

523 3.3.3. Actuators

524 Edge systems may potentially have the ability to actuate
 525 key pieces of infrastructure which can influence the environment
 526 (such as traffic lights, or barriers). Depending on what the
 527 actuator is, the Edge device(s) may be capable of having a large
 528 impact on the behaviour and security of vehicles. For example, a compromised Edge might claim a
 529 certain actuation state that is not true, such as claiming a traffic light is green when it is red.

Example Attacks
<ul style="list-style-type: none"> • Disable

530 3.3.4. Energy System

531 The energy system being used to power the Edge system is
 532 important to consider as different kinds could be used. Typically
 533 Edge systems will be powered using mains power and the attacks
 534 on this system relate to removing access to this power. However,
 535 alternate power systems (such as via batteries and renewable energy like solar) may be used in areas
 536 where providing mains power is infeasible or too costly.

Example Attacks
<ul style="list-style-type: none"> • Disconnect power supply

537 3.3.5. Physical IO

538 The Edge will have Physical IO ports that allow technicians
 539 to connect directly to the Edge infrastructure. These ports should
 540 be protected using physical security mechanisms (such as locks)
 541 to protect against attacks. From a cyber security perspective the
 542 ports need to defend against attacks that occur once physical security is bypassed. This means that any
 543 user connecting via these ports should be correctly authenticated and forensic logs made about these
 544 connection attempts.

Example Attacks
<ul style="list-style-type: none"> • Privilege Escalation

545 3.3.6. Monitoring and Logging




Figure 3. Cloud Reference Architecture (Hybrid Functional-Interaction viewpoint)

546 Records of actions taken by both the Edge and Cloud will
 547 need to be kept. This is to allow investigators to understand why
 548 a specific sequence of action occurred. They will also be needed
 549 to understand performance characteristics of the system.

550 3.3.7. Microservices

551 Microservices involve an application or services designed to
 552 provide functionality via a collection of loosely coupled services.
 553 These microservices each provide a single service compared to
 554 a monolithic model which provides multiple services at once.
 555 Benefits to this style of architecture include: improved scalability
 556 to a large number of users and increased resistance to certain attacks. A modular architecture is easier
 557 to test and develop, reducing the likelihood of bugs and vulnerabilities being present. Any services
 558 that are used internally do not need to be exposed to the wider internet, which reduces the attack
 559 surface compared to a monolithic application. However, while each individual microservice has a
 560 smaller attack surface, the inter-microservices communications become a possible avenues of attack.

561 3.3.8. Application Programming Interface (APIs)

562 The APIs exposed by a service hosted on the Edge are used
 563 to access that service. APIs can be exposed in a number of ways,
 564 however, a common technique is to use RESTful APIs [61] that
 565 represent a request and response in JSON which is typically sent
 566 over HTTP(S). As APIs often involve user provided data, it is
 567 important to ensure that it is sanitised before being manipulated or used for a task. A lack of sanitation,
 568 or vulnerabilities in the parsing code of the request can lead to confidentiality or integrity violations. A
 569 common example of this kind of attack are SQL injections.

570 3.4. Cloud Reference Architecture

571 The interactions with CAVs and the Cloud, and the operation of the Cloud are important to
 572 consider with respect to the attack surface of autonomous vehicles. Much of the information that
 573 CAVs request will be provided by Cloud services and specific applications will require interaction with
 574 Cloud APIs for services to function. The Cloud reference architecture is intended to be a simplified
 575 representation of the key components that are important for CAVs. It is sufficiently detailed for an

Example Attacks

- Delete/Modify logs
-

Example Attacks

- Malicious firmware deployment
 - Privilege Escalation
-

Example Attacks

- Lack of user data validation (e.g., SQL injection)
 - Incorrect data disclosure
-

576 analysis of how attacks on the Cloud will impact a CAV, however, more detailed reference architectures
 577 and threat models should be used to analyse the Cloud in greater depth (such as [62–64]).

578 The remainder of this section will describe the components in the Cloud reference architecture.
 579 Certain components have been previously described in [Subsection 3.3](#) (e.g., Monitoring and Logging,
 580 Microservices, APIs) and will not be repeated here.

581 3.4.1. Communication

582 The communication patterns that occur in the Cloud will
 583 be more complex due to the Cloud’s need for scalability, high
 584 performance and high reliability. Rather than having a single
 585 connection to the wider networking infrastructure, the Cloud
 586 will have multiple gateways which utilise load balancing to
 587 improve performance. As the Cloud is internet connected, large services will be under attack from
 588 DDoS packet spam [65]. This means that firewalls and DDoS protection is an important part of the
 589 Cloud’s communication infrastructure.

Example Attacks

- Jam or disconnect link
 - MiTM
 - DDoS
-

590 3.4.2. Data Storage

591 Cloud data storage will be different to both vehicular and
 592 RSU data storage, as it will be physically distributed across
 593 many different data centres. The data will also be replicated
 594 to ensure integrity and availability under hardware failures. This
 595 replication and distribution increases the attack surface of the
 596 data storage, as there are multiple sites to consider exploiting
 597 and the communication between sites to perform the replication
 598 could also be vulnerable to exploitation.

Example Attacks

- Insider attacks against data centre [66]
 - Hardware failures limiting availability
 - Unintended remote access
-

599 3.4.3. Data Analysis

600 The data analysis performed by the Cloud is going to be
 601 different from that performed by the vehicle, as the Cloud
 602 will have access to much more data over a longer time period.
 603 Therefore, the Cloud will have different objectives in terms of the
 604 analysis it produces from the data. For example, it may analyse
 605 historical data to better predict traffic patterns, which could be used to load balance road networks
 606 when a vehicle requests a route from its origin to its destination. An attacker may wish to gain this
 607 analysis (as it is likely to be very valuable) or impact the analysis so it outputs poor results (e.g., such
 608 that all vehicles are directed into a lower capacity road, leaving higher capacity roads free).

Example Attacks

- Privacy leakage of user information (Privacy Preserving Data Mining to protect it [67])
-

609 4. Methodology

610 In the previous section we presented the four components our reference architecture that can be
 611 used as an aid for the examination of cyber security threats and to develop appropriate strategies to
 612 address these threats. This reference architecture provides an abstracted view of the ecosystem, that
 613 allows developers of new products, services and infrastructure to see how their own contribution
 614 fits into this system of systems. To identify and mitigate attacks using the reference architecture, the
 615 users undertake three steps: instantiate the architecture with their particular use case; isolate the attack
 616 surface; and identify attack entry points in the boundary and internal interaction points. We explain
 617 each of these steps below.




Figure 4. The process of analysing an attack goal when performing an attack surface analysis.

618 4.1. Instantiating the Reference Architecture

619 Thus far the *abstract* reference architecture has been presented, with abstract components such
 620 as Sensors. To use the reference architecture the abstract components needs to be instantiated with
 621 concrete components as required by the specific scenario of interest (as will be shown in Section 5).
 622 For example, the Sensors component could be instantiated with multiple sensors such as LIDAR,
 623 odeometry and temperature when an application needs the output from these sensors to perform its
 624 function. Not all components need to be instantiated, as the scenario may not involve certain systems
 625 within the vehicle. Only once the reference architecture has been instantiated with components the
 626 threats against those components can be identified. Using the reference architecture the threats posed
 627 by compromised components can be identified by the links specifying how the components interact.

628 4.2. Synthesis of Attack Surface

629 Once the system has been instantiated for a use case, attack surface analysis is used to identify
 630 a comprehensive set of feasible methods for adversaries to achieve their goals. Attackers can use or
 631 combine different attack paths to reach their desired goals. Where mitigations should be implemented
 632 can be identified by focusing on reducing the ability for an adversary to exploit critical attack
 633 surfaces. Attack goals can be obtained by systematically performing a threat modelling on the
 634 critical components or functionality of a system. There are a number of approaches to perform threat
 635 modelling [6], of which Microsoft's STRIDE is commonly used in the automotive security domain. A
 636 reference architecture is useful in conjunction with threat modelling, as it provides a methodology
 637 to identify the attack routes to achieve a goal that may not have been previously considered in the
 638 threat modelling. However, performing a threat modelling is out-of-scope for this work to ensure
 639 generalisation to arbitrary threat modelling techniques.

640 One effective method to describe attack surfaces are *attack trees*, which were first introduced
 641 in [15] to manage the large number of threats derived from comprehensive threat modelling in general
 642 security. Attack tress have since been employed in automotive security in a number of scenarios [68–70].
 643 To create attack trees potential threat agents and their goals in compromising the system first need to
 644 be identified. For each attack goal, the relevant attack surfaces need to be specified that define possible
 645 paths to reach this goal. These paths can then be represented as an attack tree. At the end of this
 646 procedure, a list of attack trees which cover known goals, sub-goals and attack methods of potential
 647 threat agents are produced.

648 In this paper, we also employ attack trees to synthesise, manage, and control the attack surface.
 649 The process to perform the attack tree analysis is illustrated in Figure 4 and described below:

- 650 1. The goal(s) of the threat actor needs to be specified.
- 651 2. Using these goals, identify the component in the reference architecture that ultimately needs to
 652 be compromised for these goals to be achieved.
- 653 3. Identify the possible entry points to the system the threat actor could exploit.
- 654 4. Using the entry point(s) calculate the path(s) that an threat actor could take to reach the target
 655 component from an external interaction.
- 656 5. Considering a threat actor's capabilities, resources and presence, prune paths that the threat
 657 actor cannot exploit.

658 Evaluation of threat agents appear at both ends of an attack tree. At the beginning, goals are
659 derived from threat agents' motivations. It is assumed that threat agents will only consider goals that
660 follow from their motivations. For example, a thief has a motivation to increase their wealth, so a goal
661 is to steal physical assets rather than cause damage. Each threat will require a specific capability to be
662 carried out, such as: technique, skills, knowledge, equipment, presence, and others. Therefore, at the
663 end of the procedure, the capability of threat agents also needs to be evaluated to check if achieving
664 the goal is feasible. If achieving the goal is not feasible, then the attack tree needs to be pruned from
665 the set of attack trees generated.

666 Existing work has been performed on identifying threat actors and their capabilities, goals,
667 resources and motivations which should be used as input to this attack surface analysis. For example,
668 a comprehensive library of threat agents for general information systems was provided by Intel [71] in
669 their Threat Agent Risk Analysis (TARA) model. This library contains information of 22 threat agents
670 and their 9 common attributes. However, many of the agents are inapplicable in to CAV security. For
671 example, the TARA list was reduced to the seven most relevant agents in [7], which included: thief,
672 owner, organised crime, mechanic, hacktivist, terrorist, and foreign government.

673 *4.3. Identify Attack Entry Points at the Boundary and Internal Interaction Points*

674 Attacks against a single component can have limited impact. Therefore, is it often the case that
675 compromised components are used to aid in attacking another component, or multiple components
676 are attacked simultaneously. These attacks are more complicated and take longer to perform, but can
677 have a greater impact on the CAV. The motivations for an attacker to attack a component via another
678 compromised component can be divided into two categories: (i) escalating attacker capability, and (ii)
679 creating greater impacts. Achieving one of these categories (or both) can be obtained by sequential
680 manipulation (attacking a component from another already compromised component), concurrent
681 manipulation (attacking two components simultaneously), or a mixture of the two manipulations.

682 *4.4. Summary*

683 This section described the procedure to synthesis the attack surface of a system described using a
684 reference architecture. To provide an insight into how to apply this technique, two case studies using
685 it are explored in the next section.

686 **5. Case Studies**

687 In this section we present two different case studies to demonstrate how to use the proposed
688 reference architecture for attack surface analysis. The procedure for creating these case studies is as
689 follows:

- 690 1. Identify a scenario where cyber security is important.
- 691 2. Instantiate the reference architecture with concrete instances of components that are present in
692 the scenario.
- 693 3. Use input from threat modelling to identify the goals, motivations, capabilities and resources of
694 an adversary.
- 695 4. Use the instantiated reference architecture to build attack trees. This facilitate the cyber security
696 analysis of the scenario by identifying how an adversary will perform attacks.
- 697 5. Finally, identify the ways in which the system can be changed to mitigate the attacks.

698 *5.1. Driverless Valet Parking*

699 The first case study is the driverless valet parking example from [20], where a driver wishes to
700 leave their vehicle at a parking garage. Once the driver leaves the vehicle, they can request the vehicle
701 to autonomously park itself by collaborating with the smart parking garage. The parking garage
702 will allocate the vehicle a parking space and provide internal maps to aid the vehicle in locating its
703 allocated space. When the driver wishes to retrieve the vehicle, a signal can be sent from a smart




Figure 5. Valet Parking: Vehicle and Devices Instantiation

704 phone to request the vehicle autonomously drives back to its owner. Figures 5 and 6 show how certain
 705 components were instantiated with concrete components. In this example the Edge resources form the
 706 smart parking garage.

707 As this case study was also examined in [7], by implementing this application we will demonstrate
 708 the differences between using our reference architecture and the one proposed in [7]. This comparison
 709 will demonstrate that our architecture allows a more detailed analysis of the attack surface due to
 710 the consideration of interactions between the vehicle and the devices & peripherals, the Edge, and
 711 the Cloud. Some different components are included that are not referenced in the example in [7].
 712 These new components are in boxes with dashed lines. They indicate certain functionality that could
 713 be involved with a valet parking system and highlight different ways in which the system could be
 714 attacked that are not covered in the previous work.

715 5.1.1. Threat Identification

716 A number of threats were identified in the original example in [7] that can also identified using
 717 the reference architecture proposed.

- 718 • **Spoof GNSS on Vehicle:** GNSS signals could be spoofed to assist a thief stealing a vehicle.
- 719 • **Modify Map via Update on Cloud:** A map update is used to force the vehicle along a route to
 720 an arbitrary destination.
- 721 • **Replay Retrieval on Device:** A thief replays a recorded signal used to retrieve the vehicle.




Figure 6. Valet Parking: Edge and Cloud Instantiation

- 722 • **Blind Range Sensor on Vehicle:** An adversary seeking to cause a crash could blind the range
 723 sensor to prevent a vehicle from knowing its distance from obstacles.
 724 • **DoS Parking Allocator on Cloud:** An adversary seeking to induce a traffic jam or freeze the
 725 parking garage could DoS the allocator, preventing vehicles from requesting new spaces.

726 Using our reference architecture (see [Figure 5](#) and [Figure 6](#)), the following additional threats have been
 727 identified:

- 728 • **DoS Parking Sensor on Edge:** Cover sensor that detects a vehicle in a parking space to reduce
 729 the availability of the parking garage.
 730 • **Information Disclosure via Vulnerable APIs on Cloud:** Vulnerable APIs can potentially execute
 731 arbitrary code (such as via SQL injection attacks), allowing an adversary to remotely obtain
 732 sensitive data about the parking garage system.
 733 • **MiTM on Edge:** A device could be placed in the parking garage that mimics roadside
 734 infrastructure. If it has a high signal strength vehicles would prefer connecting to that overt V2I
 735 rather than the Edge infrastructure of the parking garage, allowing a MiTM attack between the
 736 vehicle and the cloud services. This attack could reveal sensitive information about the user
 737 (such as financial details). It could also be used to over allocate vehicles causing a large traffic
 738 jam in the parking garage (denying vehicle availability).

739 This work does not perform the threat modelling in step 3, as it is expected to input this
 740 information from one of the many different threat modelling techniques.

741 Considering the additional components contained within the dashed lines, the following additional
 742 threats have been identified:

- 743 • **Cut Mains Power on Edge:** A vehicle should be able to autonomously exit the parking garage
 744 even if power is lost, whether the power loss is malicious or not. If the vehicle is unable to exit
 745 the parking garage then availability of the vehicle is denied to its owner.

Table 3. Attack Tree Analysis for Driverless Valet Parking Use Case with example threat actors and their goals.

TA	Goal(s)	Attacked Functions	Attack Surfaces	Detailed attacks on assets
Thief	Steal the CAV	F1 Stop the CAV at location that is convenient to steal	Sensors that are responsible to stop the CAV in incidents; OR edge (can ask CAV to stop)	A12 or A22
		F2 Mislead the CAV to false location by falsifying the route	Cloud (giving false map); OR Edge (giving false location); OR GNSS sensor (responsible for location sensing)	A21 or A11 or A23 or A24
		F3 Control the CAV: compromise the command to make it go to false location	Edge (giving false command); OR Key (control the CAV directly)	A31 or A32
Hactivist	Manipulate the CAV operation	F1 Stop the CAV	(See Thief analysis)	(See Thief analysis)
		F2 Mislead the CAV	(See Thief analysis)	(See Thief analysis)
		F3 Control the CAV	(See Thief analysis)	(See Thief analysis)
		F4 Track the CAV	Cloud (storing location information of the CAV)	A41
Terrorist	G1 Manipulate the CAV operation to create accident or damage	G1-F1 Stop the CAV	(See Thief analysis)	(See Thief analysis)
		G1-F2 Mislead the CAV	(See Thief analysis)	(See Thief analysis)
		G1-F3 Control the CAV	(See Thief analysis)	(See Thief analysis)
	G2 Disrupt the station operation	G2-F5 Stop parking management services	Cloud; Or Edge	F5: A21 or A25

746 • **Incorrect Indoors Positioning on Edge:** If the Edge assists the vehicle perform indoors
747 positioning, then spoofed and jammed signals could be used to decrease the vehicle's certainty
748 of its position.

749 By including the additional interactions with the Edge and Cloud, as well as a better structuring of
750 components and their interactions, our reference architecture has allowed more threats to be identified.
751 The identification process does not require specifying a large amount of details compared to more
752 comprehensive reference architectures such as [8].

753 Attack trees were built from these specified potential threats by first selecting the most relevant
754 threat agents; which are the thief, hactivist, and terrorist. For each threat agent the most important
755 goals were identified. The attack trees were then analysed for each of these goals, which is summarised
756 in Table 3. Finally, each tree was combined into a single tree (Figure 7) to illustrate the security analysis
757 of the use case with respect to the threats, threat agents and their goals.

758 5.1.2. Discussion

759 One of the interactions specified in [7] was a key/remote that is used to initiate the retrieval of a
760 vehicle from the parking garage. This occurs by the key communicating directly with the vehicle and
761 initiating its automated driving to exit the parking garage. An alternate architecture involves a user
762 using an app on a smart phone to contact a cloud service to request the retrieval of a vehicle. This means
763 the parking garage has greater control over vehicle parking allocating and scheduling vehicle retrieval.




Figure 7. Attack Tree for Valet Parking Example that highlights multiple ways in which an attacker can reach its goal.

764 We use the alternate model as we believe it more accurately describes the way in which this system
 765 will be implemented. This means the steps of vehicle retrieval goes from (a.i) Key communicates with
 766 vehicle to request autonomous retrieval, (a.ii) vehicle autonomously drives to owner through parking
 767 garage, to (b.i) User requests vehicle retrieval using smart phone, (b.ii) cloud service schedules vehicle
 768 retrieval to prevent traffic jams in parking garage, (b.iii) vehicle autonomously drives to owner.

769 Our proposed reference architecture allows a more comprehensive attack surface to be identified
 770 than in previous work. For example, with the valet parking application, we can identify additional
 771 threats compared to [7], because we define the previous unspecified Devices & Peripherals and Cloud
 772 sub-architectures and more thoroughly define the Edge sub-architecture. A comprehensive attack
 773 surface is always important because a missed threat can lead to risk being underestimated, and may
 774 create severe consequences if attackers can use it to exploit the system.

775 5.2. Tesla Exploitation

776 An example of an attack that needed to compromise multiple components to gain further control
 777 of the vehicle was presented in [43]. The default behaviour of the vehicle was to connect to an
 778 unprotected WiFi hotspot and opening a website in the infotainment. By setting up an alternate WiFi
 779 hotspot with the same SSID but broadcasting at a higher power, the vehicle instead connected to the
 780 alternate hotspot, which allowed traffic to be redirected to a custom server. This means an attacker
 781 with *semi-local* presence could perform this attack. To perform the initial attack the adversary needed to
 782 identify a vulnerability in the web browser running in the infotainment system. By chaining together
 783 multiple browser vulnerabilities the adversary could execute arbitrary code through the compromised
 784 browser. Privilege escalation was then required to affect the system in substantial ways. Without
 785 gaining the privilege escalation the adversary would have little ability to affect the internal systems.
 786 However, read access to memory storage was provided through the browser exploit, which revealed
 787 debug information that included procedures for upgrading firmware. This allowed a custom firmware
 788 to be flashed to certain components, which is capable of performing arbitrary tasks.

Table 4. Attack Tree Analysis for Tesla Use Case with example threat actors and their goals.

TA	Goal(s)	Attacked Functions	Attack Surfaces	Detailed attacks on assets
Hacktivist	HG: to control the CAV components (e.g. IC, Parrot, Gateway) remotely	HF1 get shell access	AS-HF1: IC, Parrot, Gateway	A-HF1: A43 or A42 or A41
		HF1.1: overwrite firmware	AS-HF1.1: Linux Kernel, Browser	A-HF1.1: A22
		HF1.1.1: get firmware address	AS-HF1.1.1: Browser	A-HF1.1.1: A21
		HF1.1.2: redirect browser to fake domain	AS-HF1.1.2: Browser, WIFI	A-HF1.1.2: A11 or A12 or A13
		HF2: get root privileges	AS-HF2: Linux Kernel	A-HF2: A31
		HF2.1: disable the security app	AS-HF2.1: Linux Kernel	A-HF2.1: A32
Terrorist	TG: to create high safety impact attack by autonomous vehicle TG1: to control the CAV remotely TG2: to monitor the CAV to find environment where it can create high safety impact (e.g. involves many people)	TG1: See attacked function analysis for HG	See attack surfaces for HG	See attacks for HG
		TG2: TF-TG2: to track the CAV and its operating environment	AS-TF2: See similar analysis in valet driving example	TG2: See similar analysis in valet driving example

789 So the summarised steps are as follows, with FUNCTIONAL components and **implementations** of
790 those components formatted differently (see more details of reference model in [Figure 8](#) and [Figure 9](#)):

- 791 1. EXTERNAL COMMUNICATIONS connects to popular **WiFi** hotspot
792 2. A malicious **WiFi** hotspot spoofs the SSID with a greater signal strength
793 3. Compromise **browser** in INFOTAINMENT
794 4. Privilege escalation in **Operating System**
795 5. STORAGE access
796 6. Flash **firmware** (change STORAGE)
797 7. Custom **firmware** eavesdrops/transmits/blocks messages on the **CAN bus** (INTERNAL
798 COMMUNICATIONS)

799 Even with access to the CAN bus safety features limited the adversary's ability to perform certain
800 actions. For example, the authors attempted to open the trunk while the vehicle was in motion, but
801 this was prevented. However, the authors found a way to block certain CAN message which allowed
802 them to open the trunk, or to disable automatic locking of doors when the vehicle was moving.

803 In response to this vulnerability several controls were added to reinforce the security of Tesla
804 vehicles, including (see [Figure 10](#)):

- 805 • **C1:** Greater isolation of the Infotainment web browser from being able to interact with other
806 parts of the system
807 • **C2:** Page Table Isolation (which prevents the kernel from accessing user mode memory and thus
808 preventing the adversary from executing code in user space)
809 • **C3:** Code Signing (to prevent untrusted code execution)

810 To the best of our knowledge, this kind of attack has only been reported three times, twice by the
811 Tencent researchers in 2016 and 2017, and once by Checkoway *et al.* [54]. For each attack, the authors
812 reported the step-by-step hacking actions, while the producers quickly provided patches/fixes to the




Figure 8. Tesla Exploit: Vehicle and Devices Instantiation

813 vulnerabilities. Attack tree analysis is an effective framework to address the implications of these
 814 attacks and how effective the patches/fixes are in eliminating the risks. Firstly, it can represent and
 815 simplify the attack, while highlighting the relevant components and their relations. Secondly, the
 816 attack can be understood in greater detail by identifying alternatives in the attack surface that could be
 817 used to achieve the similar goals. Thirdly, attack implications can be drawn from considering all the
 818 related threat agents and their motivations. Finally, the controls that manufacturers added to the CAV
 819 can also be verified and the attack trees can also suggest other effective controls to be considered in
 820 mitigating the threats.

821 The security analysis from [43] is extended by our attack tree analysis in Table 4 and Figure 10,
 822 with alternatives to browser attacks being identified as part of the attack surface. The goals of the
 823 hackers (i.e. Tencent researchers) can be extended to a higher-impact goal of a terrorist, for example,
 824 to control the CAV to create accidents when it is operating in a crowded environment. Relations
 825 between the threats, goals, and agents are illustrated in Figure 10. When applying the Tesla controls,
 826 it can be seen that the detailed attacks are eliminated, the relevant attack surfaces are significantly
 827 reduced, while the connections between the surfaces are removed.

828 Finally, the ability for the manufacturer to develop a fix in a short time period is important, it is
 829 also important that the fix can be rapidly deployed to vehicles on the road. In this instance Tesla was
 830 able to create and deploy a fix for these issues in two weeks. The infrastructure support to widely




Figure 9. Tesla Exploit: Edge and Cloud Instantiation

831 deploy these firmware fixes in a short period of time allows the impact to the vehicle's occupants to
 832 be small. If the firmware update needed to be shipped to customers to install themselves, or vehicles
 833 needed to be recalled, the risk to drivers will be higher due to the longer time the vehicle's spend
 834 unpatched. Having an over-the-air update system could introduce the potential for malicious or buggy
 835 firmware to be deployed to vehicles, however, that is a trade-off that needs to be considered with
 836 respect to the ability to widely deploy an update in a short period of time.

837 6. Discussion

838 Having described the reference architecture and presented two case studies that demonstrate how
 839 to apply it, this section will now discuss some of the implications and issues raised.

840 6.1. Expectation that the CAV Architecture Will Change

841 We expect that in the next decade and beyond the functionality of a CAV will change in unexpected
 842 ways. This reference architecture is designed to reflect the functionality that is expected to be deployed
 843 in the near future. For the far future, we expect that changes will need to be made to the reference
 844 architecture, which is why it has been designed to be modular. If new components or new interactions
 845 between components need to be added, then the reference architecture can be updated to include them.
 846 In doing so the attack surface of the system will change and the analysis will need to be re-performed.

847 6.2. Prioritising Attack Surface Analysis

848 Given the limited resources, defenders need to prioritise specific attack surfaces to protect, starting
 849 with threats that pose a high risk (high likelihood and high impact). Therefore, defenders need to
 850 perform a risk assessment which takes into account threat agents' capabilities and motivations as
 851 well as the available controls in the system. However, performing a risk assessment is complicated,
 852 takes time as it involves a large number of threats, and may contain uncertainties in the calculated




Figure 10. Attack Tree for Tesla Example that highlights a sequential attack for the attacker to reach its goal.

853 risk. Moreover, CAV risks are not static due to its dynamic operating environments. Consequently,
 854 the risk assessment will need to be repeated frequently to adapt to such changes. Therefore, to make
 855 the analysis more effective, it is important to shape the focus by prioritising the security resources for
 856 several core parts.

857 We argue that the core parts that should be prioritised are critical functional components exploited
 858 frequently by threat agents. When proper controls are applied the corresponding attack surfaces will
 859 be reduced, which creates further challenges for attackers. For CAVs these components are: (i)
 860 Communication, (ii) Sensors, and (iii) Data Analysis. Wireless or physical communication is the vector
 861 through which many cyber-attacks will be perpetrated, as it acts as a gateway between external agents
 862 and the internal components. Sensors are important because they provide information about CAV's
 863 surrounding environments. If sensor information is unavailable or modified maliciously, CAV may be
 864 manipulated to make harmful decisions. Finally, the Data Analysis is important, as it influences CAV's
 865 autonomous functions.

866 6.3. The Need to Understand Trust-levels in All the Surfaces

867 Dominic *et al.* [7] recommended that the defenders should not place too much trust in individual
 868 CAV components. If any single trusted component exists, it can be the single point of failure that
 869 manipulates the whole security of the system once compromised. Consequently, defenders should
 870 put redundant security resources in different components to cross-check each other. However, when
 871 attackers manipulate more than one component, they may also be able to compromise the cross-check,
 872 eliminating a source of redundancy. Therefore, it is also important to understand trust levels properly

873 in each component. When there are inconsistencies between them, understandings of their trust-levels
874 will decide which components are in favour for making security decisions.

875 *6.4. Isolating Critical Subsystems*

876 Depending on applications and stakeholders' interest, some components can be considered more
877 critical than others. For example, safety applications emphasise more on driving functionalities, while
878 privacy applications focus more on data-related components. Putting more security resources on these
879 critical components will not be enough to secure them, given the connections between attack surfaces
880 can bring unknown threats from other vulnerable surface as shown in previous sections. Therefore, it
881 is also important to isolate these critical parts from other vulnerable surfaces, or at least to create a
882 secured shield around them by putting proper controls in their connections.

883 *6.5. Considerations of Hardware and Software Security*

884 In this reference architecture, we chose not to include the physical viewpoint and have only
885 included a virtual implementation viewpoint as full representations of both viewpoints increase the
886 difficulty of performing a high-level security analysis of a CAV. However, it is important to consider
887 the security of the hardware and software of these systems. An issue is that for vehicular systems
888 the software is typically only available as a black box, as manufactures are in general unwilling to
889 supply the source code used for implementation. The same is typically true for the hardware in a
890 CAV. This means that it can be useful to consider the system in terms of its functional components
891 and their interactions. The high-level reference architecture presented in this work will be useful to
892 initially describe the system, but a more detailed modelling language (such as SysML [72]) may be
893 preferable when more details need to be specified. However, a reference architecture will be useful
894 to highlight the attacker's path to achieve its goal and select in which component or interaction to
895 implement mitigations for the attack. Additionally, it can also be referred in security verification when
896 upgrading software or hardware for the system, which can happen frequently in a CAV's life cycle.

897 *6.6. Using Reference Architecture to Mitigate Attacks*

898 The reference architecture can provide an identification of which components and interactions
899 are critical to the attacker research a goal. By analysing the generated attack trees, the component
900 or interaction in which a mitigation is implemented can be justified. A common desire is to apply
901 security controls in all potentially vulnerable components to minimise the attack surface. However,
902 security resources are often limited, therefore, it is necessary to prioritise which countermeasures are
903 implemented. The reference architecture can help to choose which mitigations to prioritise due to the
904 ability to demonstrate the impact the mitigation will have in general. For example, the attack surfaces
905 which lead to critical impacts should have the highest priority; while restricting surfaces which open
906 the chances to attack other surfaces is usually more efficient than restricting isolated attack surfaces.

907 The reference architecture can also useful at the design phase of a system. For example, if security
908 is critical, the designers should reduce the use of components that have large attack surface (e.g., by
909 replacing them with more secure components) or restricting access to insecure functionality that link
910 to other critical functionality. Finally, in the long term, the reference architecture can help to manage
911 the complexity of systems and attacks. For instance, it can be used to visualise new vulnerabilities at a
912 high level and also identify relevant mitigations when the system design changes.

913 **7. Future Work**

914 In this section two key areas in which the attack surface analysis needs to be developed further
915 are discussed: (i) the automation of the analysis of system, and (ii) how to understand dynamically
916 changing risk in different environments and scenarios.

917 7.1. Automated Analysis

918 In this work the reference architecture and the attack surface analysis has been performed
919 manually. The important components, their interactions, and the ways in which they can be attacked
920 have been derived by analysing how CAVs operate and how they can be attacked. Alternatively,
921 the reference architecture and the attack surface analysis could be automated. To achieve this the
922 important components and their interactions would need to be manually identified. These interactions
923 could be specified in terms of the kind of interaction they represent (for example the class of data that is
924 sent from one component to another). With this information the attack surface could be automatically
925 explored using information about how an adversary could attack the system. This would allow attack
926 trees to be automatically generated. However, not all attacks are likely to be interesting or feasible, so
927 some manual pruning would be required.

928 7.2. Understanding Dynamic Risk

929 Risk analysis has been becoming compulsory to understand and control the potential system
930 breaches and vulnerabilities [69]. Moreover, this analysis can also be used to rank the threats to help
931 defenders deploy security resource most effectively for a mitigation plan. Although extensive research
932 has been carried out on CAV risk analysis, there is little study which adequately tackle dynamic risks
933 that CAVs are facing. It is not sufficient to assess CAV risks just for a single time because as a moving
934 system, CAV's environment is changing frequently. As a result, risk assessments need updating to
935 reflect new knowledge of environments and systems [73].

936 In the future, we plan to investigate factors that affect CAV risk assessment by answering when
937 and how new assessments are needed, and the most efficient way to manage dynamic risks. This
938 research would be essential to help CAVs to adapt quickly and more appropriately when operating in
939 dynamic environments.

940 8. Conclusion

941 In this paper we have presented a reference architecture from a hybrid functional-communication
942 viewpoint. This combined viewpoint allows easier attack surface analysis as the components and their
943 interactions can be analysed from a single diagram. This reference architecture has been designed
944 with four key sub-architectures for CAVs, the Edge, the Cloud, and Devices & Peripherals. The latter
945 three are key to understanding the attack surface of a CAV, because they present new attack vectors
946 that have previously been hard to specify. Finally, two examples of how to instantiate the reference
947 architecture and analyse that instantiation have been provided showing how new and existing attacks
948 can be analysed using this reference architecture.

949 **Author Contributions:** Conceptualization, Carsten Maple; Methodology, Matthew Bradbury, Kevin Ghirardello
950 and Anh Tuan Le; Formal Analysis, Anh Tuan Le; Writing—Original Draft Preparation, Matthew Bradbury, Kevin
951 Ghirardello and Anh Tuan Le; Writing—Review & Editing, Matthew Bradbury, Kevin Ghirardello, Anh Tuan
952 Le, and Carsten Maple; Visualization, Matthew Bradbury and Anh Tuan Le; Supervision, Carsten Maple; Project
953 Administration, Carsten Maple; Funding Acquisition, Carsten Maple

954 **Funding:** This work is supported by the Alan Turing Institute under EPSRC grant EP/N510129/1, the UK Hub
955 for Cyber Security of the Internet of Things, PETRAS, under grant (EP/N02334X/1) and the CAPRI project under
956 grant (TS/P012264/1).

957 **Acknowledgments:** The authors would like to thank Daniel Fowler for assistance proofreading this work.

958 **Conflicts of Interest:** The authors declare no conflict of interest. The founding sponsors had no role in the design
959 of the study; in the collection, analyses, or interpretation of data; in the writing of the manuscript, and in the
960 decision to publish the result.

961 Abbreviations

962 The following abbreviations are used in this manuscript:

963

AD	Autonomous Driving
CAN	Controller Area Network
CAV	Connected and Autonomous Vehicles
DoS	Denial of Service
DDoS	Distributed Denial of Service
DSRC	Dedicated Short-Range Communications
ECU	Electronic Control Unit
GNSS	Global Navigation Satellite System (such as GPS, GLONASS, Galileo and BeiDou)
ITS	Intelligent Transport Systems
LIDAR	Light Detection and Ranging (detects object distance using light)
MiTM	Man-in-the-Middle
RCE	Remote Code Execution
RSU	Roadside Unit
TaaS	Theft as a Service
UAV	Unmanned Air Vehicle
USV	Unmanned Sea Vehicle

964 References

- 965 1. Winners of £51 million government competition to develop world-leading self-driving car testing
 966 infrastructure unveiled. [https://www.gov.uk/government/news/winners-of-51-million-government-](https://www.gov.uk/government/news/winners-of-51-million-government-competition-to-develop-world-leading-self-driving-car-testing-infrastructure-unveiled)
 967 [competition-to-develop-world-leading-self-driving-car-testing-infrastructure-unveiled](https://www.gov.uk/government/news/winners-of-51-million-government-competition-to-develop-world-leading-self-driving-car-testing-infrastructure-unveiled), 2017. Accessed:
 968 2018-07-31.
- 969 2. Siegel, J.E.; Erb, D.C.; Sarma, S.E. A Survey of the Connected Vehicle Landscape—Architectures, Enabling
 970 Technologies, Applications, and Development Areas. *IEEE Transactions on Intelligent Transportation Systems*
 971 **2018**, *19*, 2391–2406. doi:10.1109/TITS.2017.2749459.
- 972 3. Hussain, R.; Zeadally, S. Autonomous Cars: Research Results, Issues, and Future Challenges. *IEEE*
 973 *Communications Surveys Tutorials* **2019**, *21*, 1275–1313. doi:10.1109/COMST.2018.2869360.
- 974 4. Hegde, R.; Mishra, G.; Gurumurthy, K.S. An Insight into the Hardware and Software Complexity of ECUs
 975 in Vehicles. *Advances in Computing and Information Technology*; Wyld, D.C.; Wozniak, M.; Chaki, N.;
 976 Meghanathan, N.; Nagamalai, D., Eds.; Springer Berlin Heidelberg: Berlin, Heidelberg, 2011; pp. 99–106.
- 977 5. The Institution of Engineering and Technology.; the Knowledge Transfer Network. *Automotive Cyber*
 978 *Security: An IET/KTN Thought Leadership Review of risk perspectives for connected vehicles*. Technical
 979 report, The Institution of Engineering and Technology, 2015.
- 980 6. SAE. *J3061: Cybersecurity Guidebook for Cyber-Physical Vehicle Systems*, 2016. J3061_201601.
- 981 7. Dominic, D.; Chhawri, S.; Eustice, R.M.; Ma, D.; Weimerskirch, A. Risk Assessment for Cooperative
 982 Automated Driving. *Proceedings of the 2nd ACM Workshop on Cyber-Physical Systems Security and*
 983 *Privacy*; ACM: New York, NY, USA, 2016; CPS-SPC '16, pp. 47–58. doi:10.1145/2994487.2994499.
- 984 8. The Architecture Team. *Architecture Reference for Cooperative and Intelligent Transportation*. [https://](https://local.iteris.com/arc-it/index.html)
 985 local.iteris.com/arc-it/index.html, 2018. Version 8.1.
- 986 9. Sheehan, B.; Murphy, F.; Mullins, M.; Ryan, C. Connected and autonomous vehicles: A cyber-risk
 987 classification framework. *Transportation Research Part A: Policy and Practice* **2019**, *124*, 523 – 536.
 988 doi:10.1016/j.tra.2018.06.033.
- 989 10. Petit, J.; Stottelaar, B.; Feiri, M. Remote Attacks on Automated Vehicles Sensors : Experiments on Camera
 990 and LiDAR. *Black Hat Europe*, 2015.
- 991 11. Miller, C.; Valasek, C. Remote Exploitation of an Unaltered Passenger Vehicle. *Black Hat*, 2015.
- 992 12. Parkinson, S.; Ward, P.; Wilson, K.; Miller, J. Cyber Threats Facing Autonomous and Connected
 993 Vehicles: Future Challenges. *IEEE Transactions on Intelligent Transportation Systems* **2017**, *18*, 2898–2915.
 994 doi:10.1109/TITS.2017.2665968.
- 995 13. Macher, G.; Höller, A.; Sporer, H.; Armengaud, E.; Kreiner, C. A Combined Safety-Hazards and
 996 Security-Threat Analysis Method for Automotive Systems. *Computer Safety, Reliability, and Security*;
 997 Koornneef, F.; van Gulijk, C., Eds.; Springer International Publishing: Cham, 2015; pp. 237–250.
- 998 14. Luettel, T.; Himmelsbach, M.; Wuensche, H. Autonomous Ground Vehicles — Concepts And A Path To
 999 The Future. *Proceedings of the IEEE* **2012**, *100*, 1831–1839. doi:10.1109/JPROC.2012.2189803.

- 1000 15. Schneier, B. Secret and Lies. *Digital Security in a Networked World* 2000.
- 1001 16. Alberts, C.J.; Dorofee, A. *Managing Information Security Risks: The Octave Approach*; Addison-Wesley
1002 Longman Publishing Co., Inc.: Boston, MA, USA, 2002.
- 1003 17. UcedaVelez, T.; Morana, M.M. *Risk Centric Threat Modeling: process for attack simulation and threat analysis*;
1004 John Wiley & Sons, 2015.
- 1005 18. McCarthy, C.; Harnett, K.; Carter, A. Characterization of potential security threats in modern automobiles:
1006 A composite modeling approach. Technical report, Washington, DC: National Highway Traffic Safety
1007 Administration, 2014. Report No. DOT HS 812 074.
- 1008 19. Taxonomy and Definitions for Terms Related to On-Road Motor Vehicle Automated Driving Systems.
1009 https://www.sae.org/standards/content/j3016_201806, 2018.
- 1010 20. Bartels, A.; Eberle, U.; Knapp, A. Deliverable D2.1: System Classification And Glossary.
1011 [http://www.adaptive-ip.eu/index.php/AdaptIVe-SP2-v12-DL-D2.1-System%20Classification-](http://www.adaptive-ip.eu/index.php/AdaptIVe-SP2-v12-DL-D2.1-System%20Classification-file=files-adaptive-content-downloads-Deliverables%20&%20papers-AdaptIVe-SP2-v12-DL-D2.1-System%20Classification.pdf)
1012 [file=files-adaptive-content-downloads-Deliverables%20&%20papers-AdaptIVe-SP2-v12-DL-D2.1-](http://www.adaptive-ip.eu/index.php/AdaptIVe-SP2-v12-DL-D2.1-System%20Classification-file=files-adaptive-content-downloads-Deliverables%20&%20papers-AdaptIVe-SP2-v12-DL-D2.1-System%20Classification.pdf)
1013 [System%20Classification.pdf](http://www.adaptive-ip.eu/index.php/AdaptIVe-SP2-v12-DL-D2.1-System%20Classification-file=files-adaptive-content-downloads-Deliverables%20&%20papers-AdaptIVe-SP2-v12-DL-D2.1-System%20Classification.pdf), 2015.
- 1014 21. Koscher, K.; Czeskis, A.; Roesner, F.; Patel, S.; Kohno, T.; Checkoway, S.; McCoy, D.; Kantor, B.; Anderson,
1015 D.; Shacham, H.; Savage, S. Experimental Security Analysis of a Modern Automobile. 2010 IEEE
1016 Symposium on Security and Privacy, 2010, pp. 447–462. doi:10.1109/SP.2010.34.
- 1017 22. Foster, I.; Prudhomme, A.; Koscher, K.; Savage, S. Fast and Vulnerable: A Story of Telematic Failures. 9th
1018 USENIX Workshop on Offensive Technologies (WOOT 15); USENIX Association: Washington, D.C., 2015.
- 1019 23. Petit, J.; Shladover, S.E. Potential Cyberattacks On Automated Vehicles. *IEEE Transactions on Intelligent*
1020 *Transportation Systems* 2015, 16, 546–556. doi:10.1109/TITS.2014.2342271.
- 1021 24. Cloutier, R.; Muller, G.; Verma, D.; Nilchiani, R.; Hole, E.; Bone, M. The Concept of Reference Architectures.
1022 *Systems Engineering* 2010, 13, 14–27, [<https://onlinelibrary.wiley.com/doi/pdf/10.1002/sys.20129>].
1023 doi:10.1002/sys.20129.
- 1024 25. Reference Architecture For Space Data Systems. <https://public.ccsds.org/Pubs/311x0m1.pdf>, 2008.
1025 CCSDS 311.0-M-1.
- 1026 26. Lin, S.W.; Miller, B.; Durand, J.; Bleakley, G.; Chigani, A.; Martin, R.; Murphy, B.; Crawford, M. The
1027 Industrial Internet of Things Volume G1: Reference Architecture. [http://www.iiconsortium.org/IIC_](http://www.iiconsortium.org/IIC_PUB_G1_V1.80_2017-01-31.pdf)
1028 [PUB_G1_V1.80_2017-01-31.pdf](http://www.iiconsortium.org/IIC_PUB_G1_V1.80_2017-01-31.pdf), 2017.
- 1029 27. CEN-CENELEC-ETSI Smart Grid Coordination Group. Smart Grid Reference Architecture. https://ec.europa.eu/energy/sites/ener/files/documents/xpert_group1_reference_architecture.pdf, 2012.
- 1030 28. Behere, S.; Törngren, M. A functional reference architecture for autonomous driving. *Information and*
1031 *Software Technology* 2016, 73, 136 – 150. doi:10.1016/j.infsof.2015.12.008.
- 1032 29. Osório, A.L.; Afsarmanesh, H.; Camarinha-Matos, L.M. Towards A Reference Architecture For A
1033 Collaborative Intelligent Transport System Infrastructure. Collaborative Networks For A Sustainable
1034 World; Camarinha-Matos, L.M.; Boucher, X.; Afsarmanesh, H., Eds.; Springer Berlin Heidelberg: Berlin,
1035 Heidelberg, 2010; pp. 469–477.
- 1036 30. Stevens, A.; Dianati, M.; Katsaros, K.; Han, C.; Fallah, S.; Maple, C.; McCullough, F.; Mouzakitis, A.
1037 Cooperative automation through the cloud: The CARMA project. 12th ITS European Congress, 2017.
- 1038 31. Sheik, A.T.; Maple, C. Key Security Challenges for Cloud-Assisted Connected and Autonomous Vehicles.
1039 Living in the Internet of Things: Cybersecurity of the IoT, 2019, pp. 1–10.
- 1040 32. Passchier, I.; van Sambeek, M. Architecture For C-ITS Applications In The Netherlands. [http://](http://smartmobilitycommunity.eu/sites/default/files/AI_CITSArchitectureNL_v1.00.pdf)
1041 smartmobilitycommunity.eu/sites/default/files/AI_CITSArchitectureNL_v1.00.pdf, 2016.
- 1042 33. Heise, C.D., Architecture Reference Of ITS In The USA. In *Intelligent Transport Systems*; Wiley-Blackwell,
1043 2015; chapter 2, pp. 18–35. doi:10.1002/9781118894774.ch2.
- 1044 34. Begoña, M.; Sergio, C.; (Iñaki), O.I.; Isabel, T.A., Reference ITS Architectures In Europe. In *Intelligent*
1045 *Transport Systems*; Wiley-Blackwell, 2015; chapter 1, pp. 1–17. doi:10.1002/9781118894774.ch1.
- 1046 35. Puñal, O.; Aguiar, A.; Gross, J. In Vanets We Trust?: Characterizing Rf Jamming In Vehicular Networks.
1047 Proceedings Of The Ninth Acm International Workshop On Vehicular Inter-networking, Systems, And
1048 Applications; ACM: New York, NY, USA, 2012; VANET '12, pp. 83–92. doi:10.1145/2307888.2307903.
- 1049 36. Davis, A. Broadcasting Your Attack: Security Testing DAB Radio In Cars. Black Hat, 2015.
- 1050

- 1051 37. Sampigethaya, K.; Li, M.; Huang, L.; Poovendran, R. AMOEBA: Robust Location Privacy
1052 Scheme for VANET. *IEEE Journal on Selected Areas in Communications* **2007**, *25*, 1569–1589.
1053 doi:10.1109/JSAC.2007.071007.
- 1054 38. Gu, P.; Khatoun, R.; Begriche, Y.; Serhrouchni, A. Vehicle Driving Pattern Based Sybil Attack Detection.
1055 2016 IEEE 18th International Conference On High Performance Computing And Communications; Ieee
1056 14th International Conference On Smart City; Ieee 2nd International Conference On Data Science And
1057 Systems (hpcc/smartcity/dss), 2016, pp. 1282–1288. doi:10.1109/HPCC-SmartCity-DSS.2016.0182.
- 1058 39. Xu, W.; Wegner, M.; Wolf, L.; Kapitza, R. Byzantine Agreement Service For Cooperative Wireless Embedded
1059 Systems. 2017 47th Annual Ieee/ifip International Conference On Dependable Systems And Networks
1060 Workshops (dsn-w), 2017, pp. 10–15. doi:10.1109/DSN-W.2017.45.
- 1061 40. Francillon, A.; Danev, B.; Capkun, S. Relay Attacks On Passive Keyless Entry And Start Systems In Modern
1062 Cars. Network And Distributed Systems Security (NDSS) Symposium, 2011.
- 1063 41. USB Kill. <https://usbkill.com>. Accessed: 2018-06-21.
- 1064 42. Currie, R. Hacking The CAN Bus: Basic Manipulation Of A Modern Automobile Through CAN Bus
1065 Reverse Engineering. Technical report, 2017.
- 1066 43. Nie, S.; Liu, L.; Du, Y. Free-fall: Hacking Tesla From Wireless To CAN Bus. Black Hat, 2017.
- 1067 44. Evenchick, E. Hopping on the CAN Bus. Black Hat Asia, 2015.
- 1068 45. Munir, A.; Koushanfar, F. Design And Analysis Of Secure And Dependable Automotive Cps: A
1069 Steer-by-wire Case Study. *IEEE Transactions on Dependable and Secure Computing* **2018**, pp. 1–1.
1070 doi:10.1109/TDSC.2018.2846741.
- 1071 46. Rouf, I.; Miller, R.; Mustafa, H.; Taylor, T.; Oh, S.; Xu, W.; Gruteser, M.; Trappe, W.; Seskar, I. Security
1072 And Privacy Vulnerabilities Of In-car Wireless Networks: A Tire Pressure Monitoring System Case Study.
1073 Proceedings Of The 19th Usenix Conference On Security; USENIX Association: Berkeley, CA, USA, 2010;
1074 USENIX Security'10, pp. 21–21.
- 1075 47. Papernot, N.; McDaniel, P.; Sinha, A.; Wellman, M. SoK: Security And Privacy In Machine Learning. 3rd
1076 IEEE European Symposium On Security And Privacy, 2018.
- 1077 48. GPS.gov. GPS Standard Positioning Service (SPS) Performance Standard. Technical Report 4, 2008.
- 1078 49. Athalye, A.; Engstrom, L.; Ilyas, A.; Kwok, K. Synthesizing Robust Adversarial Examples. 35th
1079 International Conference On Machine Learning (icml 2018), 2018. To Appear.
- 1080 50. Golson, J. Tesla driver killed in crash with Autopilot active, NHTSA investigating. [https://](https://www.theverge.com/2016/6/30/12072408/tesla-autopilot-car-crash-death-autonomous-model-s)
1081 www.theverge.com/2016/6/30/12072408/tesla-autopilot-car-crash-death-autonomous-model-s, 2016.
1082 Accessed: 2018-07-31.
- 1083 51. Sasiadek, J.Z.; Hartana, P. Sensor data fusion using Kalman filter. Proceedings of the Third
1084 International Conference on Information Fusion, 2000, Vol. 2, pp. WED5/19–WED5/25 vol.2.
1085 doi:10.1109/IFIC.2000.859866.
- 1086 52. Ivanov, R.; Pajic, M.; Lee, I. Attack-resilient Sensor Fusion For Safety-critical Cyber-physical Systems. *ACM*
1087 *Trans. Embed. Comput. Syst.* **2016**, *15*, 21:1–21:24. doi:10.1145/2847418.
- 1088 53. Davidson, D.; Wu, H.; Jellinek, R.; Singh, V.; Ristenpart, T. Controlling UAVs With Sensor Input Spoofing
1089 Attacks. 10th USENIX Workshop On Offensive Technologies (WOOT 16); USENIX Association: Austin,
1090 TX, 2016.
- 1091 54. Checkoway, S.; McCoy, D.; Kantor, B.; Anderson, D.; Shacham, H.; Savage, S.; Koscher, K.; Czeskis, A.;
1092 Roesner, F.; Kohno, T. Comprehensive Experimental Analyses of Automotive Attack Surfaces. Proceedings
1093 of the 20th USENIX Conference on Security; USENIX Association: Berkeley, CA, USA, 2011; SEC'11, pp.
1094 6–6.
- 1095 55. Woo, S.; Jo, H.J.; Lee, D.H. A Practical Wireless Attack On The Connected Car And Security
1096 Protocol For In-vehicle Can. *IEEE Transactions on Intelligent Transportation Systems* **2015**, *16*, 993–1006.
1097 doi:10.1109/TITS.2014.2351612.
- 1098 56. Kikuchi, H.; Yokomizo, T. Location Privacy Vulnerable From Bluetooth Devices. 2013 16th International
1099 Conference On Network-based Information Systems, 2013, pp. 534–538. doi:10.1109/NBiS.2013.89.
- 1100 57. Li, Z.; Pei, Q.; Markwood, I.; Liu, Y.; Pan, M.; Li, H. Location Privacy Violation Via Gps-agnostic
1101 Smart Phone Car Tracking. *IEEE Transactions on Vehicular Technology* **2018**, *67*, 5042–5053.
1102 doi:10.1109/TVT.2018.2800123.

- 1103 58. Volkswagen Apps. <http://www.volkswagen.co.uk/about-us/innovation/mobile-apps>. Accessed:
1104 2018-06-21.
- 1105 59. Cunche, M. I Know Your Mac Address: Targeted Tracking Of Individual Using Wi-fi. *Journal of Computer*
1106 *Virology and Hacking Techniques* **2014**, *10*, 219–227. doi:10.1007/s11416-013-0196-1.
- 1107 60. Gañán, C.; Loo, J.; Ghosh, A.; Esparza, O.; Reñé, S.; Muñoz, J.L. Analysis Of Inter-rsu Beaconing Interference
1108 In Vanets. Multiple Access Communications; Bellalta, B.; Vinel, A.; Jonsson, M.; Barcelo, J.; Maslennikov,
1109 R.; Chatzimisios, P.; Malone, D., Eds.; Springer Berlin Heidelberg: Berlin, Heidelberg, 2012; pp. 49–59.
- 1110 61. Fielding, R.T. Architectural Styles and the Design of Network-based Software Architectures. PhD thesis,
1111 2000. AAI9980887.
- 1112 62. Liu, F.; Tong, J.; Mao, J.; Bohn, R.; Messina, J.; Badger, L.; Leaf, D. NIST cloud computing reference
1113 architecture. *NIST special publication* **2011**, *500*, 1–28.
- 1114 63. Juliadotter, N.V.; Choo, K.R. Cloud Attack and Risk Assessment Taxonomy. *IEEE Cloud Computing* **2015**,
1115 *2*, 14–20. doi:10.1109/MCC.2015.2.
- 1116 64. Alhebaishi, N.; Wang, L.; Jajodia, S.; Singhal, A. Threat Modeling For Cloud Data Center Infrastructures.
1117 Foundations And Practice Of Security; Cuppens, F.; Wang, L.; Cuppens-Boulahia, N.; Tawbi, N.;
1118 Garcia-Alfaro, J., Eds.; Springer International Publishing: Cham, 2017; pp. 302–319.
- 1119 65. Mansfield-Devine, S. The growth and evolution of DDoS. *Network Security* **2015**, *2015*, 13 – 20.
1120 doi:[https://doi.org/10.1016/S1353-4858\(15\)30092-1](https://doi.org/10.1016/S1353-4858(15)30092-1).
- 1121 66. Duncan, A.J.; Creese, S.; Goldsmith, M. Insider Attacks In Cloud Computing. 2012 IEEE 11th International
1122 Conference On Trust, Security And Privacy In Computing And Communications, 2012, pp. 857–862.
1123 doi:10.1109/TrustCom.2012.188.
- 1124 67. Bertino, E.; Lin, D.; Jiang, W., A Survey of Quantification of Privacy Preserving Data Mining Algorithms. In
1125 *Privacy-Preserving Data Mining: Models and Algorithms*; Aggarwal, C.C.; Yu, P.S., Eds.; Springer US: Boston,
1126 MA, 2008; pp. 183–205. doi:10.1007/978-0-387-70992-5_8.
- 1127 68. Henniger, O.; Apvrille, L.; Fuchs, A.; Roudier, Y.; Ruddle, A.; Weyl, B. Security requirements for automotive
1128 on-board networks. Intelligent Transport Systems Telecommunications,(ITST), 2009 9th International
1129 Conference on. IEEE, 2009, pp. 641–646.
- 1130 69. Boudguiga, A.; Boulanger, A.; Chiron, P.; Klaudel, W.; Labiod, H.; Seguy, J.C. RACE: Risk analysis for
1131 cooperative engines. New Technologies, Mobility and Security (NTMS), 2015 7th International Conference
1132 on. IEEE, 2015, pp. 1–5.
- 1133 70. Monteuuis, J.P.; Boudguiga, A.; Zhang, J.; Labiod, H.; Serval, A.; Urien, P. SARA: Security Automotive Risk
1134 Analysis Method. Proceedings of the 4th ACM Workshop on Cyber-Physical System Security. ACM, 2018,
1135 pp. 3–14.
- 1136 71. Rosenquist, M. Prioritizing information security risks with threat agent risk assessment. *Intel Corporation*
1137 *White Paper* **2009**.
- 1138 72. Object Management Group. *OMG Systems Modeling Language*, 2017. Version 1.5.
- 1139 73. Zio, E. The future of risk assessment. *Reliability Engineering & System Safety* **2018**, *177*, 176–190.
1140 doi:10.1016/j.res.2018.04.020.