Lancaster University

Towards More Effective Performance Fuzzing

Nov. 2022

Yiqun Chen, Matthew Bradbury and Neeraj Suri

y.chen101@lancaster.ac.uk

Agenda

Lancaster University

- Motivation
- Background
- Early experiments

Motivation

Lancaster University

- Performance issues
- Worst algorithmic cases
  - Example: Bad map (dictionary) implementation
  - Constant complexity -> linear complexity
- Bad algorithm -> bad performance
- DoS attack: from *< 1 min* to *44 min*

Motivation

Lancaster University

- Many approaches for performance diagnostics: mostly profilers
- How to determine performance issues in the first place?
- Fuzzing: automatic generation of test cases

Background – Fuzzing

Lancaster University

- A random process
  - Run test cases with inputs
  - Trace coverage
  - Select inputs (guide the fuzzing)
  - Mutate inputs
  - Repeat
- Searching for inputs yield larger code coverage

Background – Performance Fuzzing

Lancaster University

Algorithm 1 The PERFFUZZ algorithm

Inputs: program p, set of inputs *Seeds*

```
1:  P ← Seeds
2:  t ← 0
3:  repeat                                              ▷ begin a cycle
4:    for input in P do
5:      with probability FUZZPROB(input) do
6:        for 1 ≤ i ≤ NUMCHILDREN(p, input) do
7:          child ← MUTATE(input)
8:          feedback ← RUN(p, child)
9:          if NEWCOV(feedback) ∨ NEWMAX(feedback) then
10:           P ← P ∪ {child}
11:         t ← t + 1
12: until given time budget expires
```

By Lemieux et al.

Motivation – Hypotheses

Lancaster University

- Fuzzing parameters
  - Test case size: 1MB
  - Timeout: 1 second
- Path length as the measure of performance fuzzing?
  - Each piece of the path (basic block) has different performance

Early Experiments

Lancaster University

- S1: default parameters (1 second 1 MiB)
- S2: custom parameter (10 seconds 100 MiB)
- Run fuzzing for 16 hours with 8 fuzzers (8 hours on 1 fuzzer on perffuzz paper)
- Measure execution time of all normal test cases and timeout test cases
- Repeat 100 times

Early Experiments

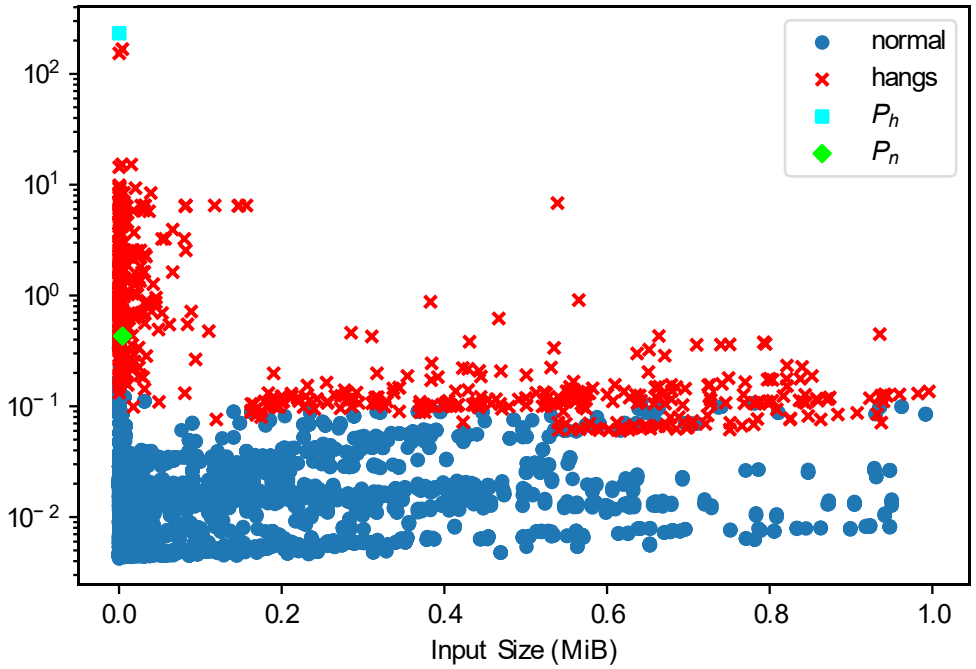

- X-axis: Test case size (MiB)
- Y-axis: Execution time (seconds)
- Blue dots: median of 100 execution times from normal test cases
- Red crosses: median of 100 execution times from timeout test cases
- Green diamond: slowest normal test case
- Teal square: slowest timeout test case

Early Experiments – File Size

Lancaster University

S1: 1 MB / 1 second (default)

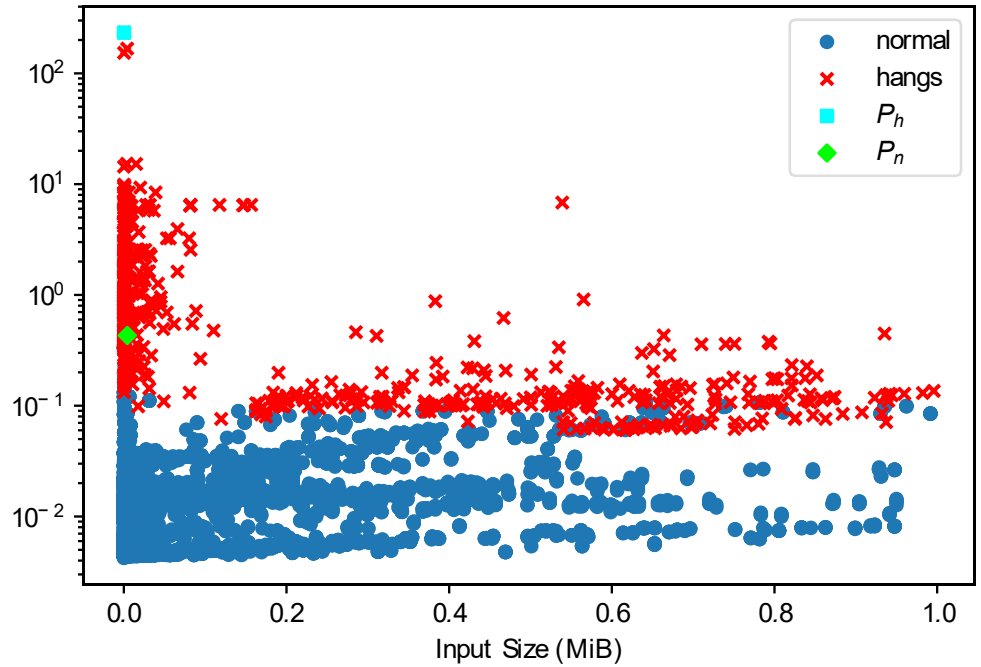

S2: 100 MB / 10 seconds

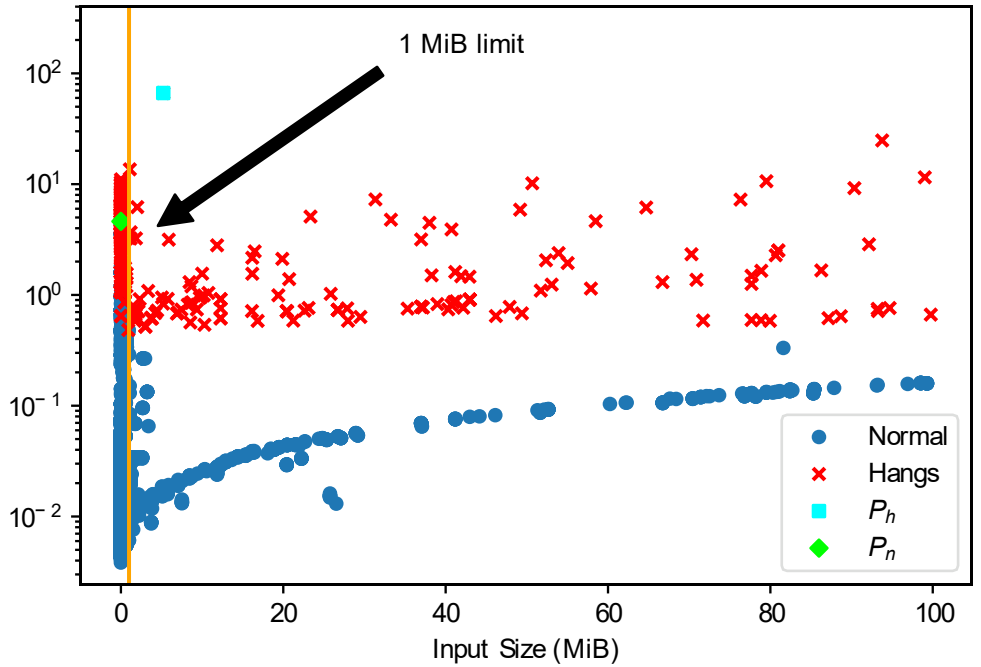

Early Experiments – Path Length

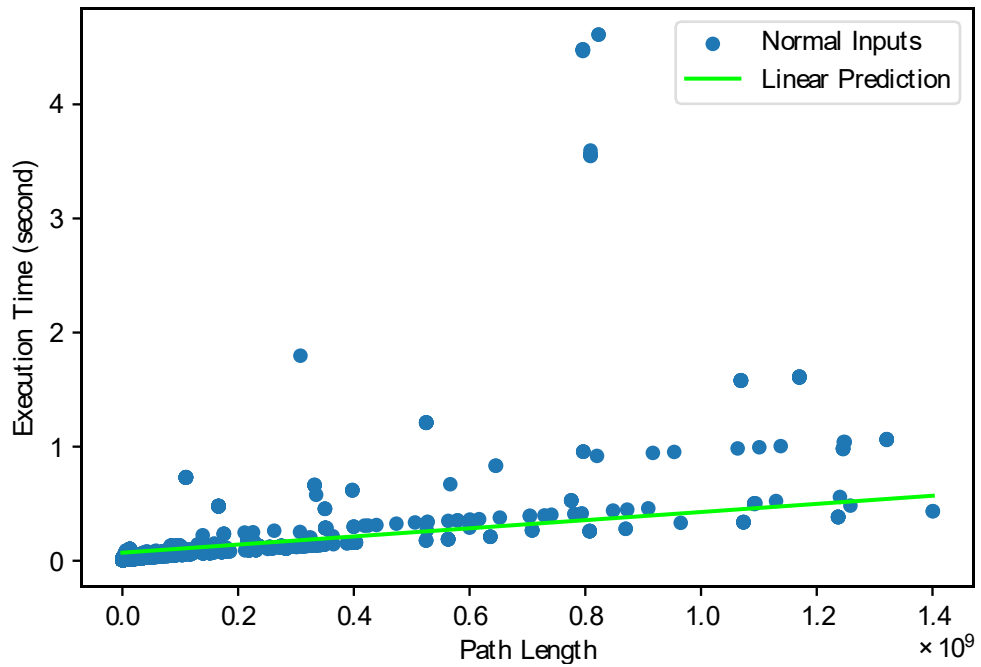

- Blue dots: median of 100 execution times
- Green line: linear regression
- Performance is somehow correlated to path length
- Non-correlated test cases are more interesting

Lancaster University

Thank you for attending, any questions?